

深度学习

Lecture 8 Auto-Regressive Models

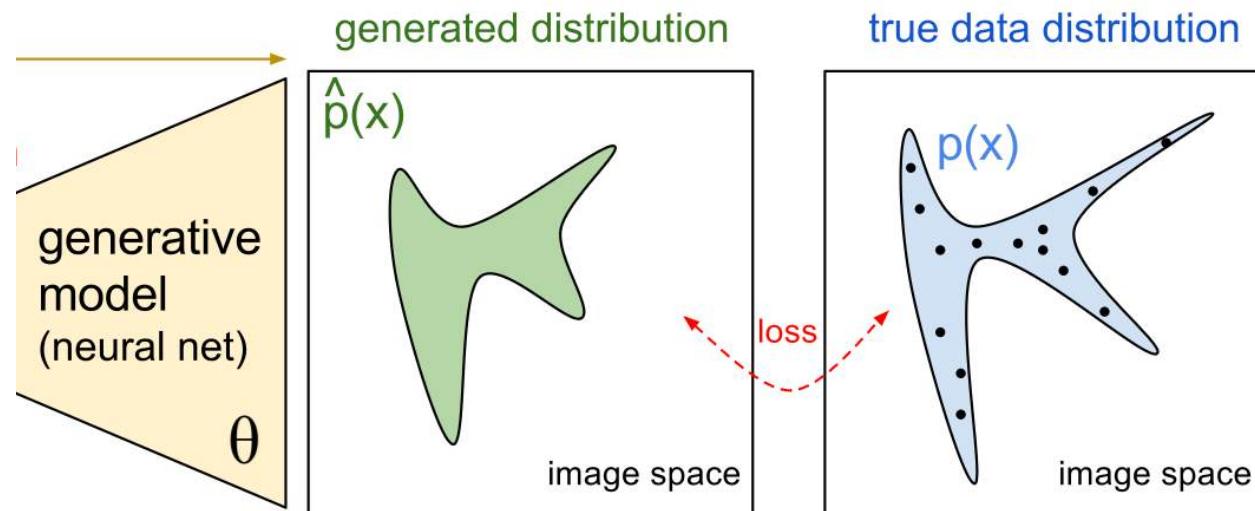
Pang Tongyao, YMSC

Generative Models



Generative Models

- Generative Models: Given some data, train a model to generate data like it!



“What I cannot create, I do not understand.” ----Richard Feynman

Generative Models

Why generative models:

- The discriminative model is to approximate the MMSE Estimator (i.e., posterior mean):

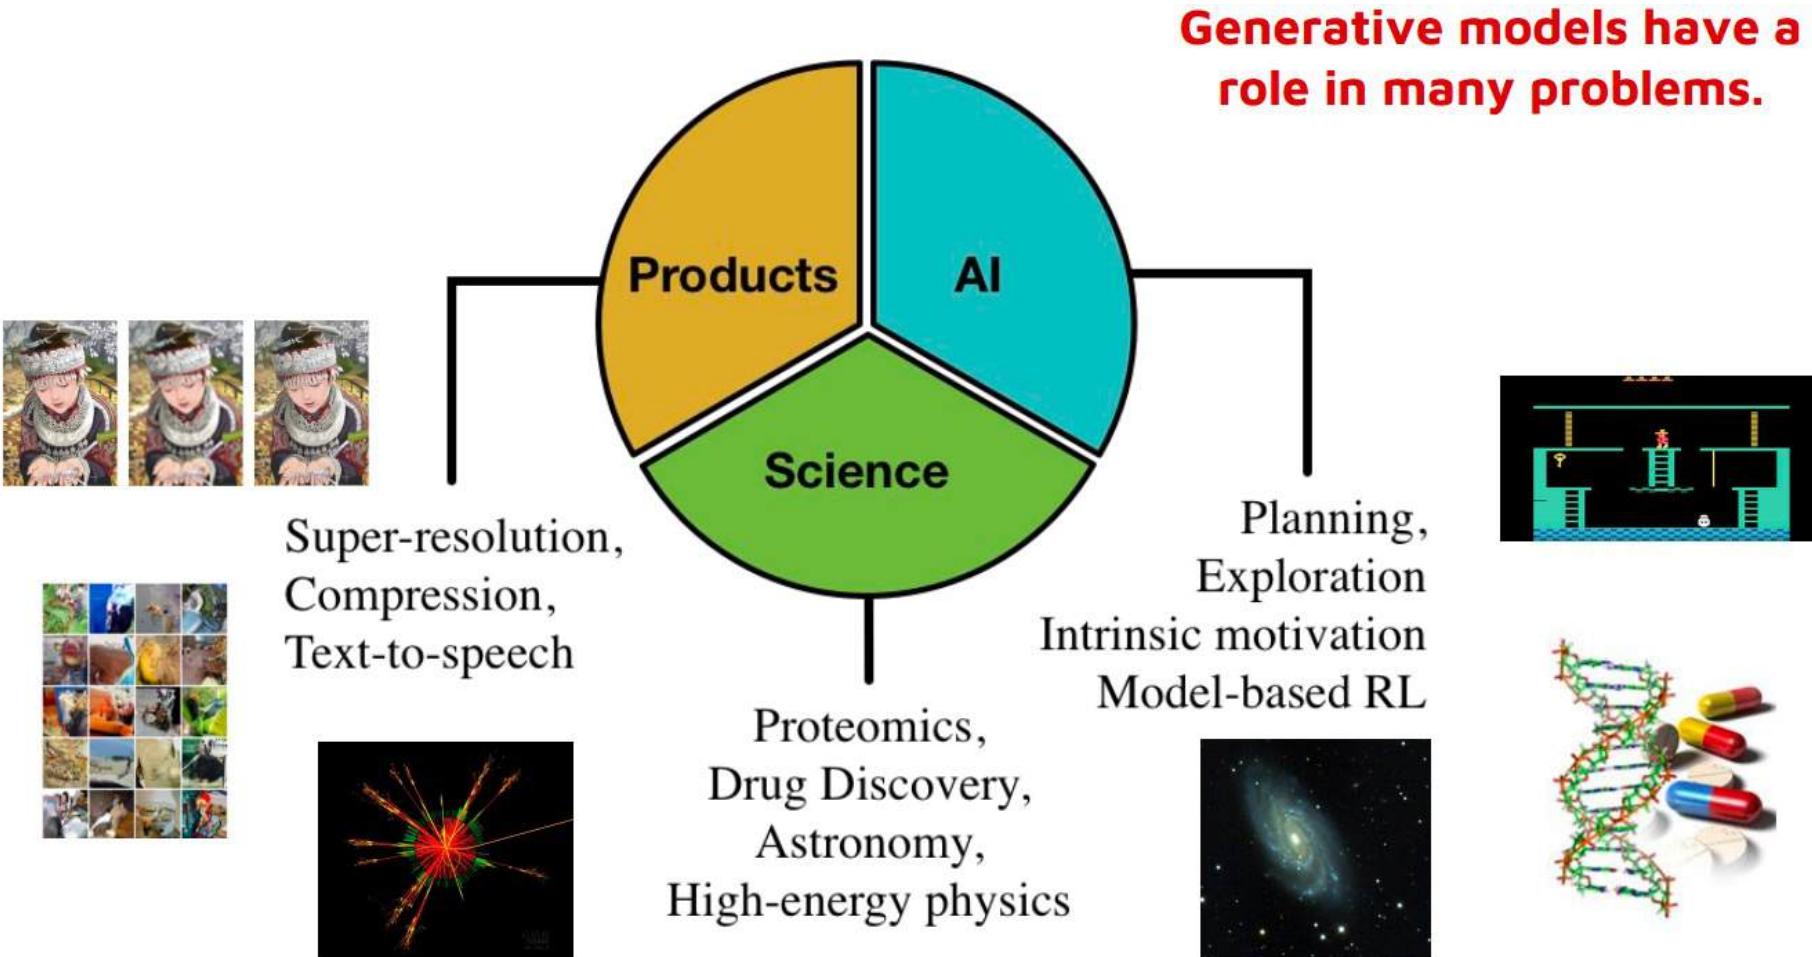
$$x_0^* = \arg \min_{g(y_0): g \in \{f_\theta\}} \mathbb{E}_{(x_0, y_0) \sim p_{data}} \|g(y_0) - x_0\|_2^2 \approx \mathbb{E}(x|y)$$

- Generative models learn $p(x)$. For any task, knowing the conditional distribution $p(y|x)$ allows us to obtain the posterior distribution

$$p(x|y) \text{ through Bayes' theorem: } p(x|y) = \frac{p(x)p(y|x)}{p(y)}.$$

- Generative models can learn the intrinsic characteristics of the data.

Generative Models



Generative Models

- Target: p_{model} is close to p_{data}
- KL divergence measures the similarities of distributions

$$KL(p_{data} \parallel p_{model}) = \int p_{data} \log p_{model} - \int p_{data} \log p_{data}$$

$$\sum_{x \in D} \log p_{model}(x|\theta)$$

loss: likelihood

↑
constant

Key: how to model and compute p_{model} ?

Autoregressive Models

Autoregressive Models

- Data in generation tasks: long text, high resolution images, codes et al.
- Model the data distribution $p(\mathbf{x})$ ($\mathbf{x} \in \mathbb{R}^N$) in an autoregressive way:

$$p_{model}(\mathbf{x}) = \prod_{i=1}^N p_{model}(x_i | x_1, \dots, x_{i-1})$$

- Autoregression is a way of modeling joint distribution by a product of conditional distributions.
- This formulation is always valid w/o compromise or approximation.

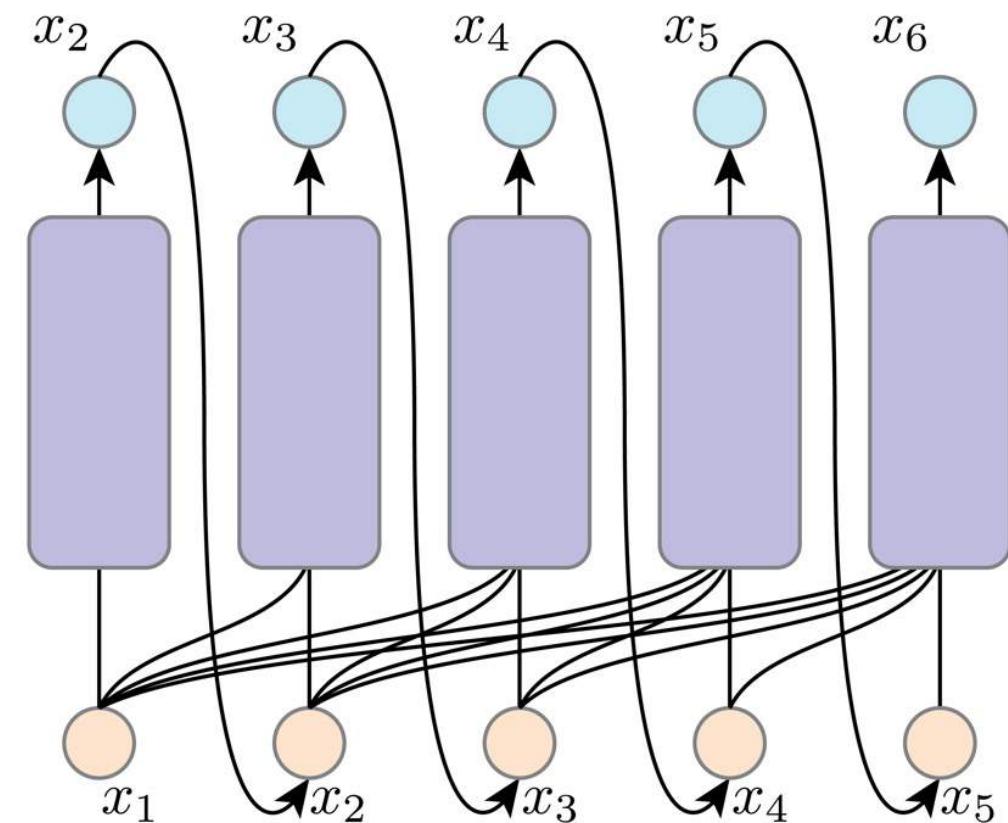
Inductive Bias

- Generally, the distributions $p(x_i|x_1, \dots, x_{i-1})$ are different for different i .
- Representing them using similar architectures helps reducing model size and improving generalization.
 - e.g. model them with shared architectures and weights (like what we did in CNN, RNN and transformers): $p_\theta(x_i|x_1, \dots, x_{i-1})$
 - In auto-regressive language models, $p_\theta(x_i|x_1, \dots, x_{i-1})$ is discrete. The network works like classification: predict the probabilities of each word in the vocabulary that appears as next word.

Autoregressive Inference

- **Auto:** using its “own” out puts as inputs for next predictions
- **Regression:** estimating relationship between variables

$$p_{\theta}(x) = \prod_{i=1}^N p_{\theta}(x_i|x_1, \dots, x_{i-1})$$



Autoregressive Inference

- This **net** models $p(x_2 | x_1)$
- **1 input**
- **1 output**

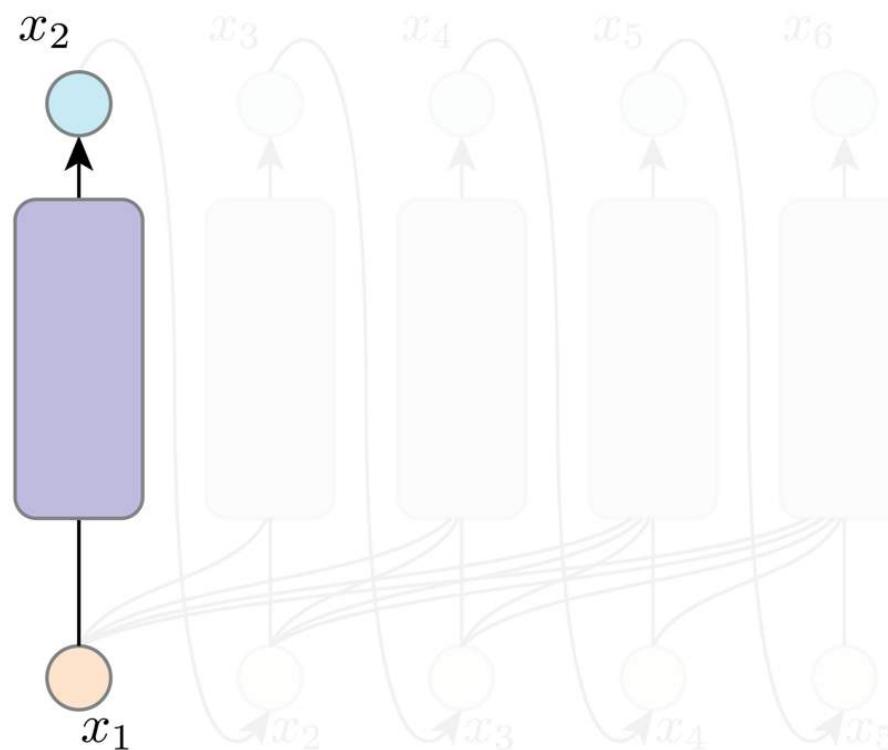
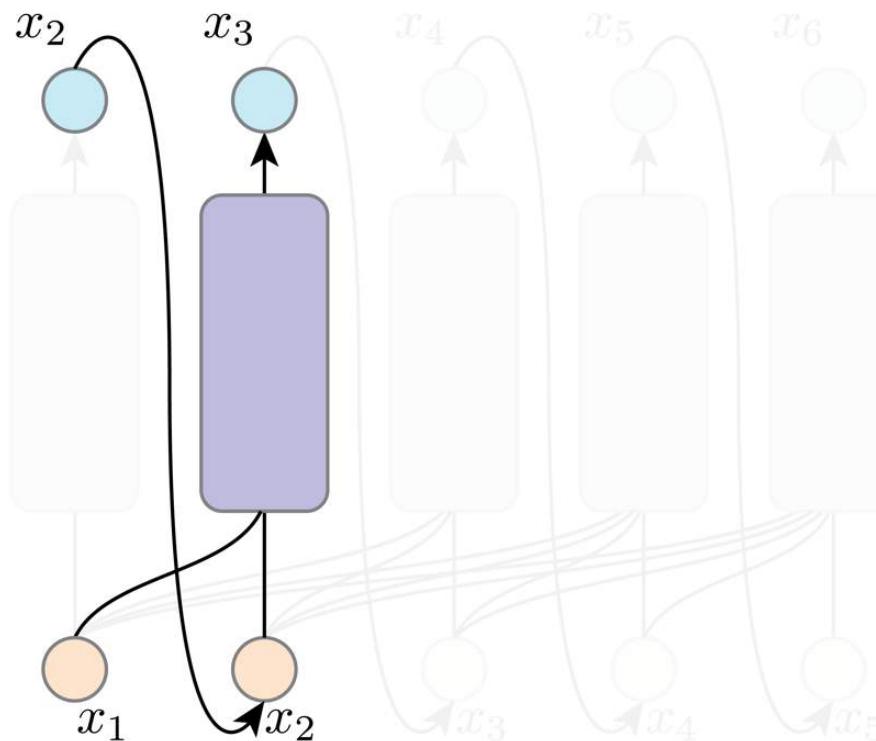


Image credit: Kaiming He

Autoregressive Inference

- This **net** models $p(x_3 | x_{1,2})$
- **2 inputs**
- **1 output**
- inputs: outputs from previous steps



Autoregressive Inference

- This **net** models $p(x_6 | x_{1,2,3,4,5})$
- **5 inputs**
- **1 output**
- inputs: outputs from previous steps

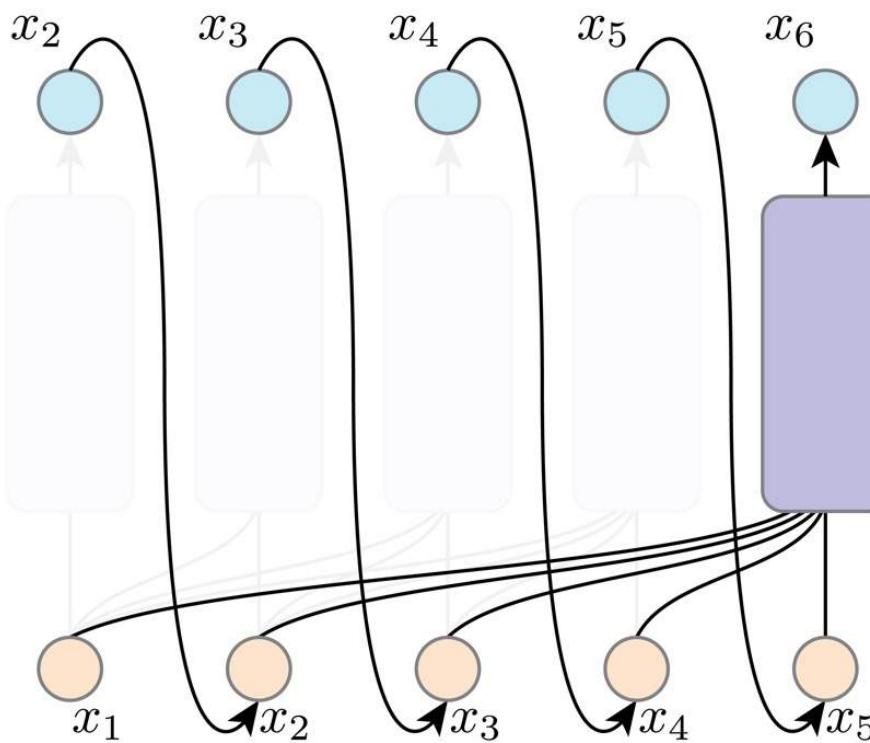


Image credit: Kaiming He

Training

- For training phase, performing the forward process in an auto-regressive way is impractical.
- Computing the gradient of x_6 should go through
 - all previous outputs
 - all previous sampling
 - all previous networks

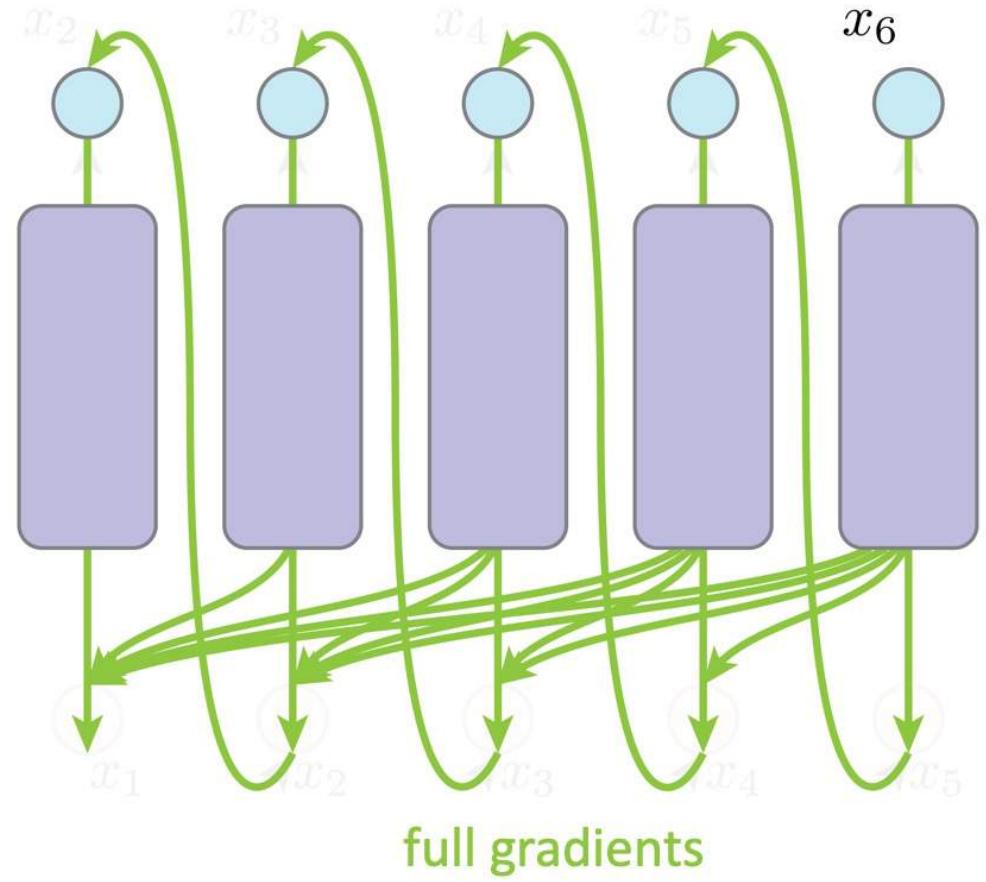


Image credit: Kaiming He

Teacher Forcing

- During **training**, we always feed the correct x_{t-1} , not the model's own prediction.
- Pros:
 - Allow parallel computation over all timesteps
 - Avoid error accumulation during training
 - Efficiently compute gradients
- Cons: train-test mismatch

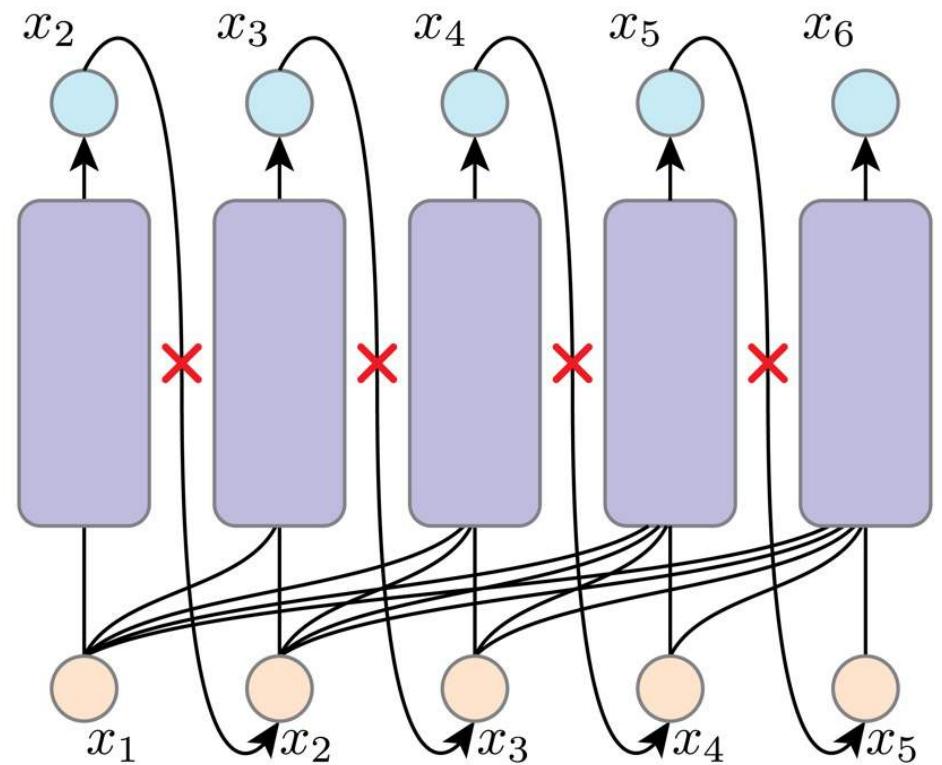
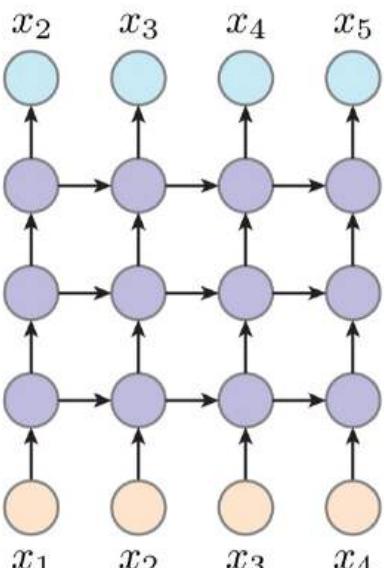
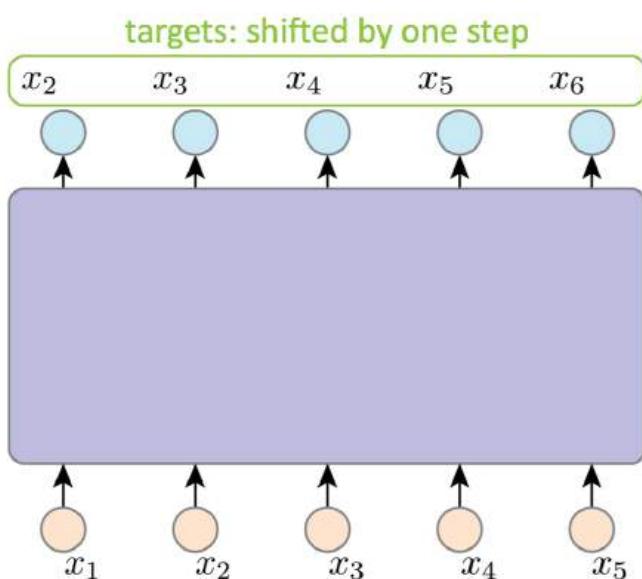


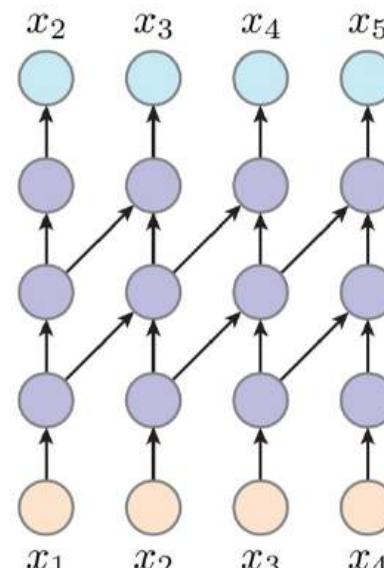
Image credit: Kaiming He

Architectures



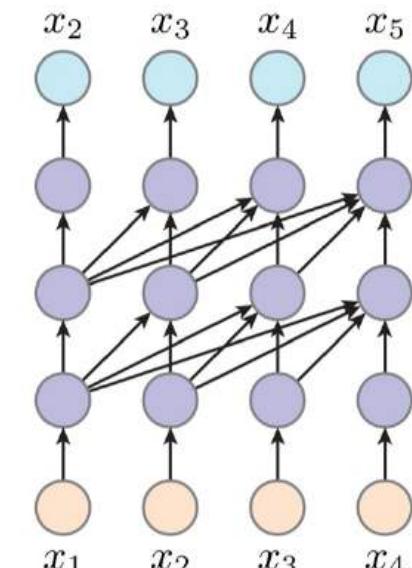
RNN

not parallelable



CNN

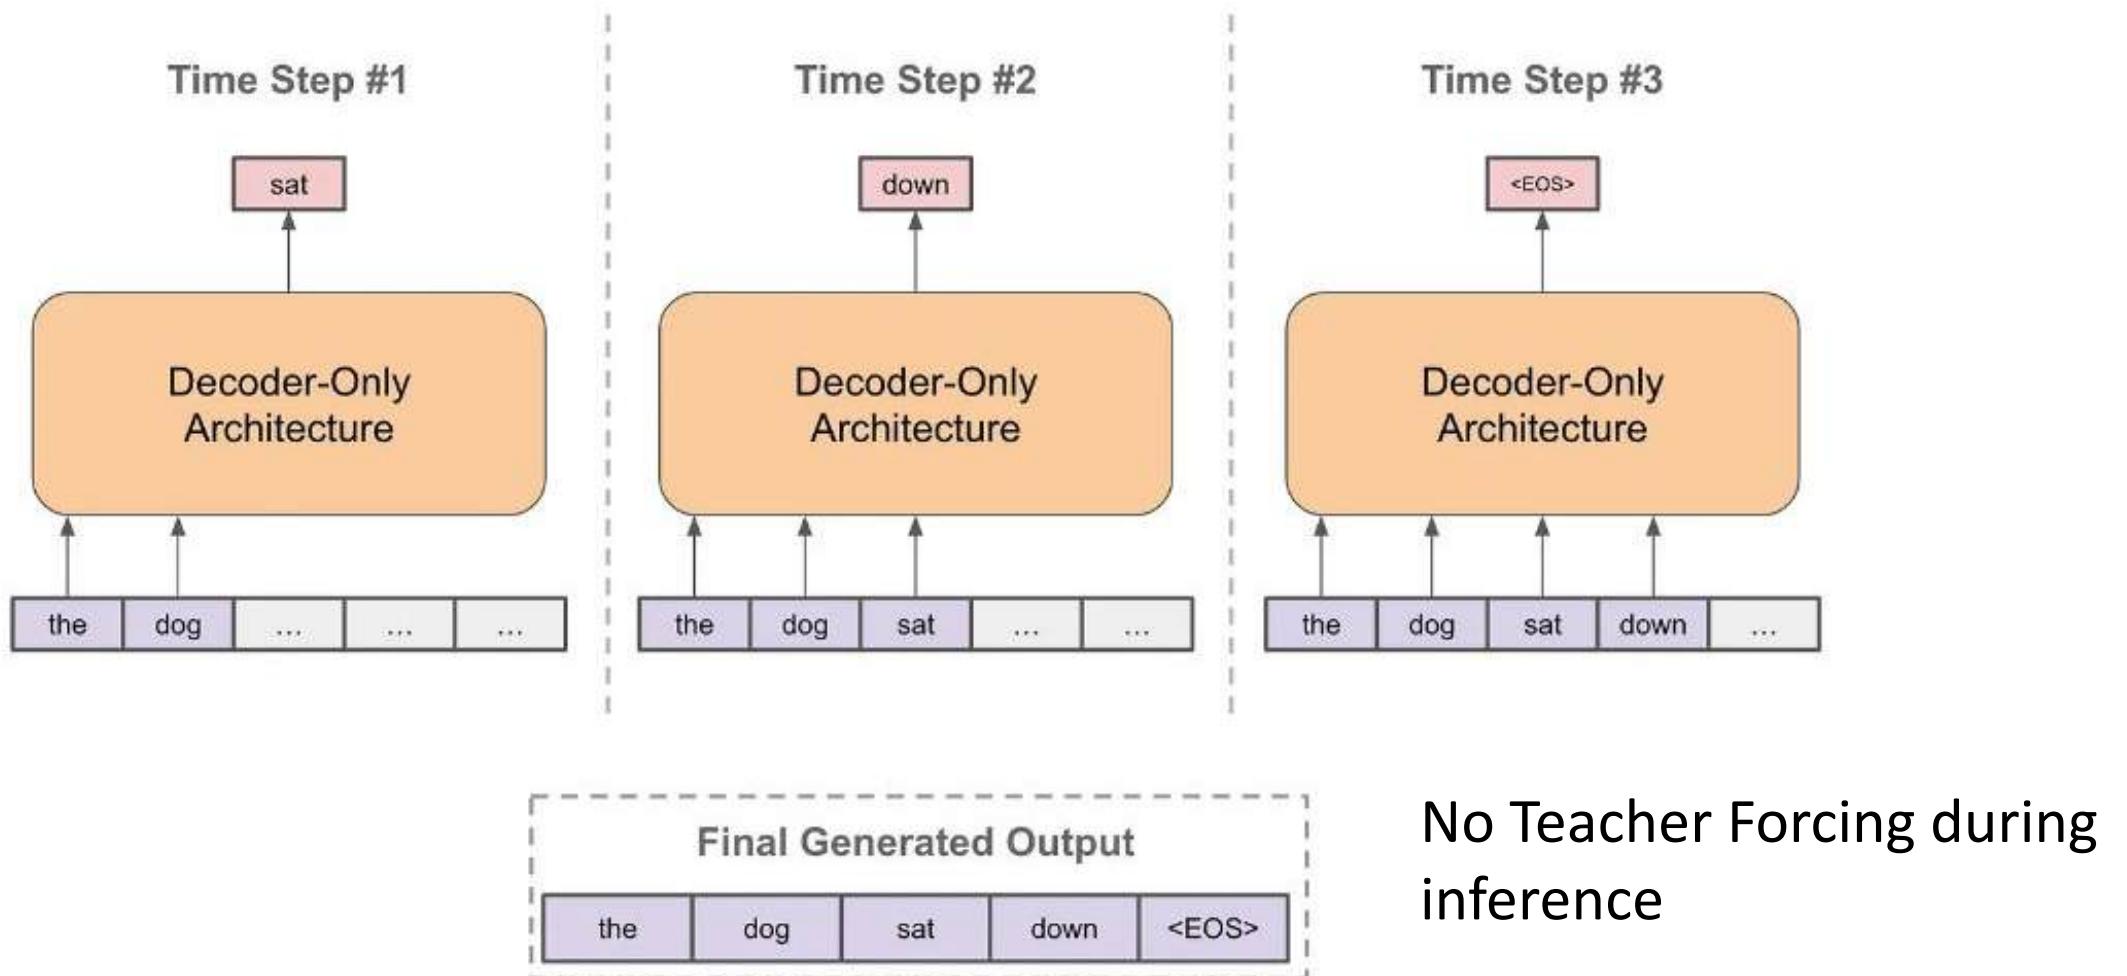
short context



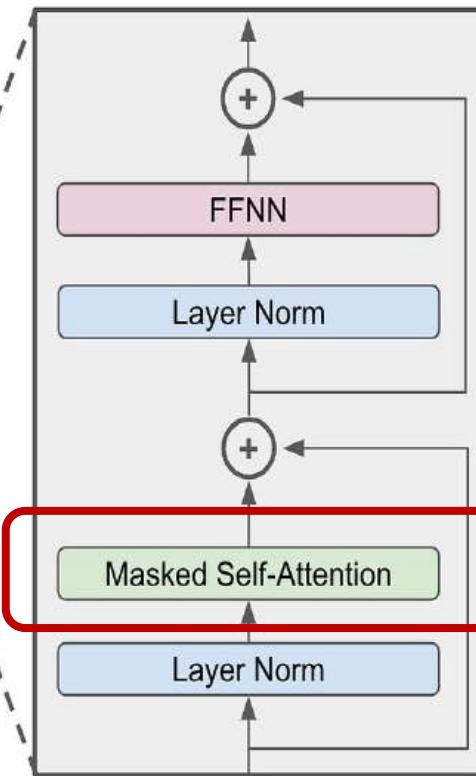
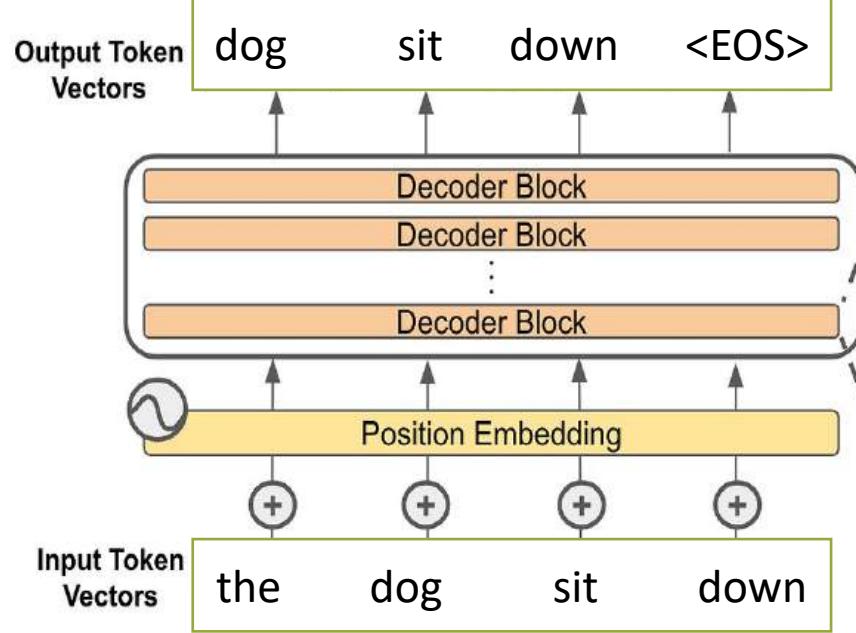
Attention

- parallelable
- long context

Autoregressive Generation



Causal Self-Attention



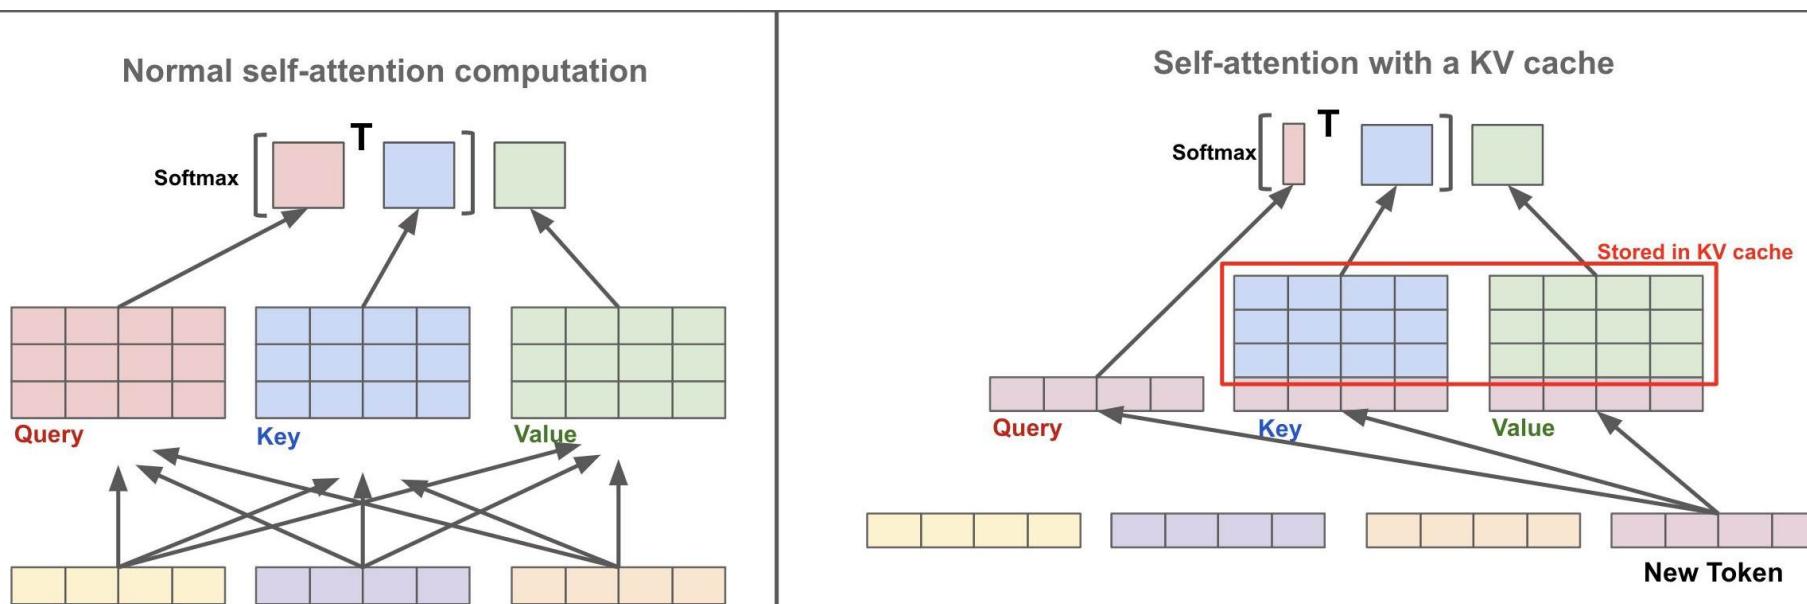
GPT: Generative Pre-Training

- **Causal Self-Attention:** Each token can only attend to itself and previous tokens — not any future tokens.

	x_1	x_2	x_3	x_4
x_1	✓			
x_2	✓	✓		
x_3	✓	✓	✓	
x_4	✓	✓	✓	✓

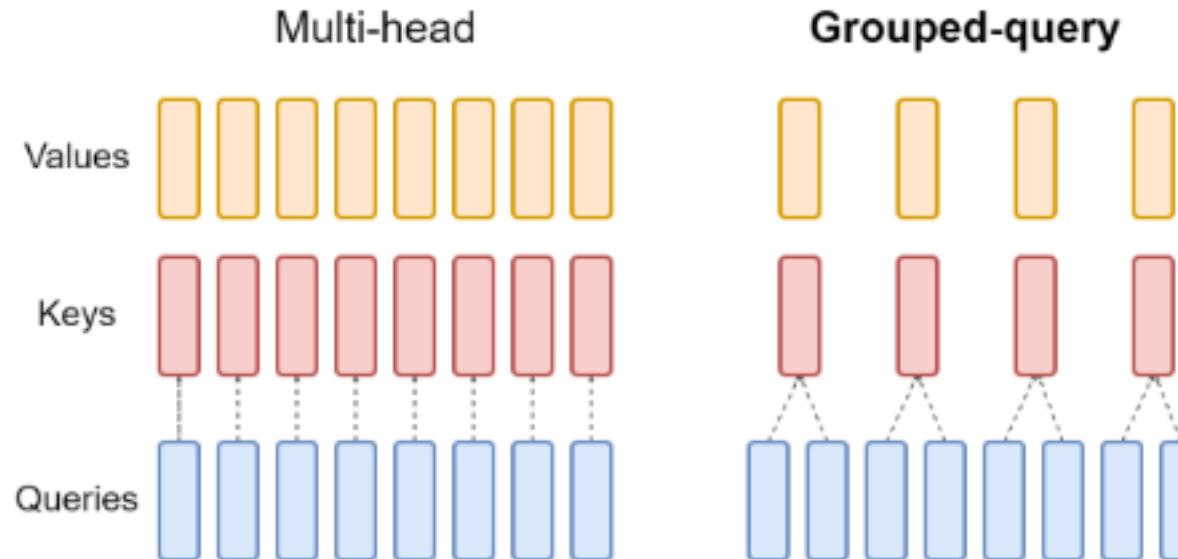
KV cache

- We only need to compute the representation for each new token, but other tokens remain fixed. KV cache size: about $2 * \text{head_dim} * \text{n_heads} * \text{n_layers} * \text{seq_length} * \text{batch_size}$



Grouped-query Attention

- For the query vectors, the original number of attention heads is maintained.
- The key and value vector pairs are then shared across multiple query heads.



Sliding Window Attention

- SWA makes the `seq_length` term a constant, allowing us to scale to larger context windows while keeping the KV cache size fixed.

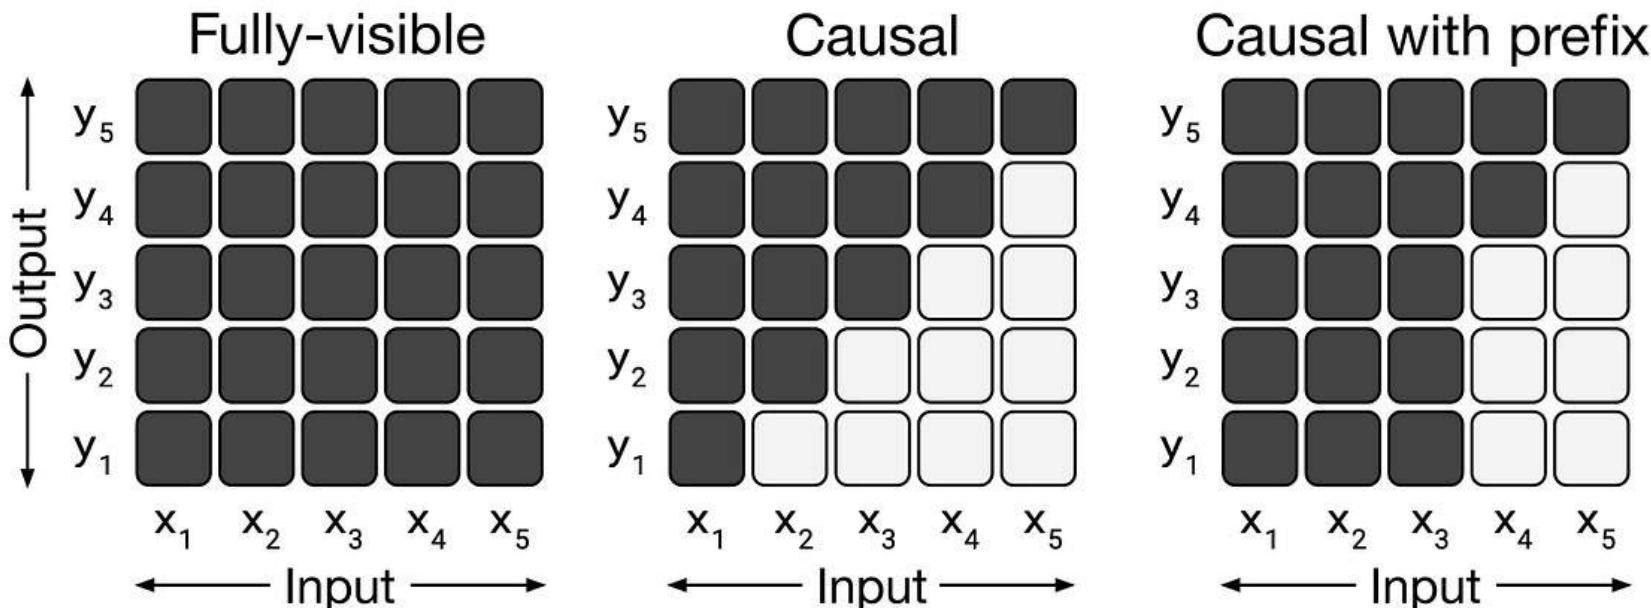
Causal Mask in Self-Attention

	The	baby	was	asleep	by	ten
The	1.00	-	-	-	-	-
baby	0.40	0.60	-	-	-	-
was	0.10	0.50	0.40	-	-	-
asleep	0.17	0.23	0.55	0.05	-	-
by	0.05	0.05	0.12	0.73	0.05	-
ten	0.04	0.08	0.03	0.15	0.17	0.53

Sliding Window Attention ($w = 3$)

	The	baby	was	asleep	by	ten
The	1.00	-	-	-	-	-
baby	0.40	0.60	-	-	-	-
was	0.10	0.50	0.40	-	-	-
asleep	-	0.23	0.55	0.05	-	-
by	-	-	0.12	0.73	0.05	-
ten	-	-	-	0.15	0.17	0.53

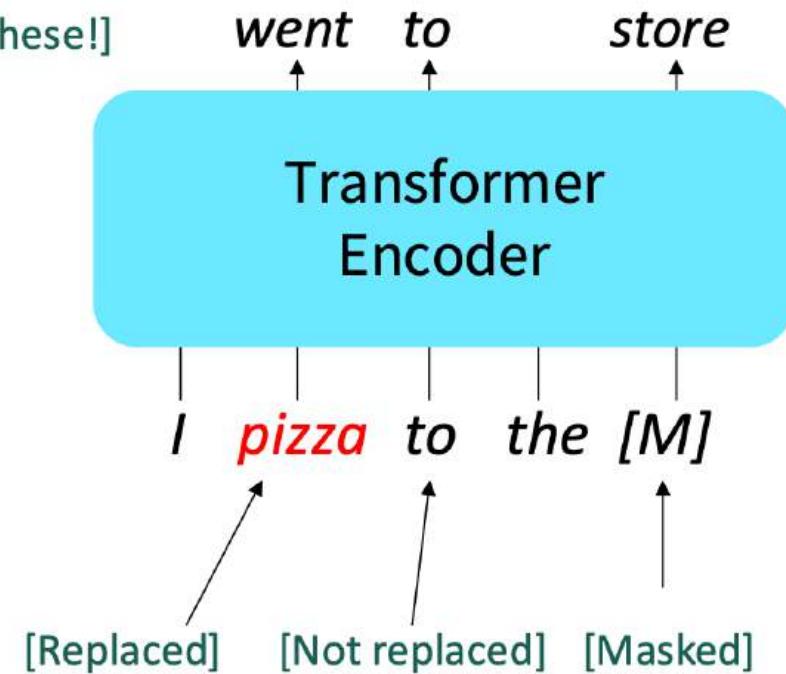
Attention Mask Patterns



- Encoder-only models leverage **bidirectional (or fully-visible)** self-attention, which considers all tokens within the entire sequence during self-attention.
- Decoder-only models use **causal self-attention**, where each token only considers tokens that come before it in the sequence.

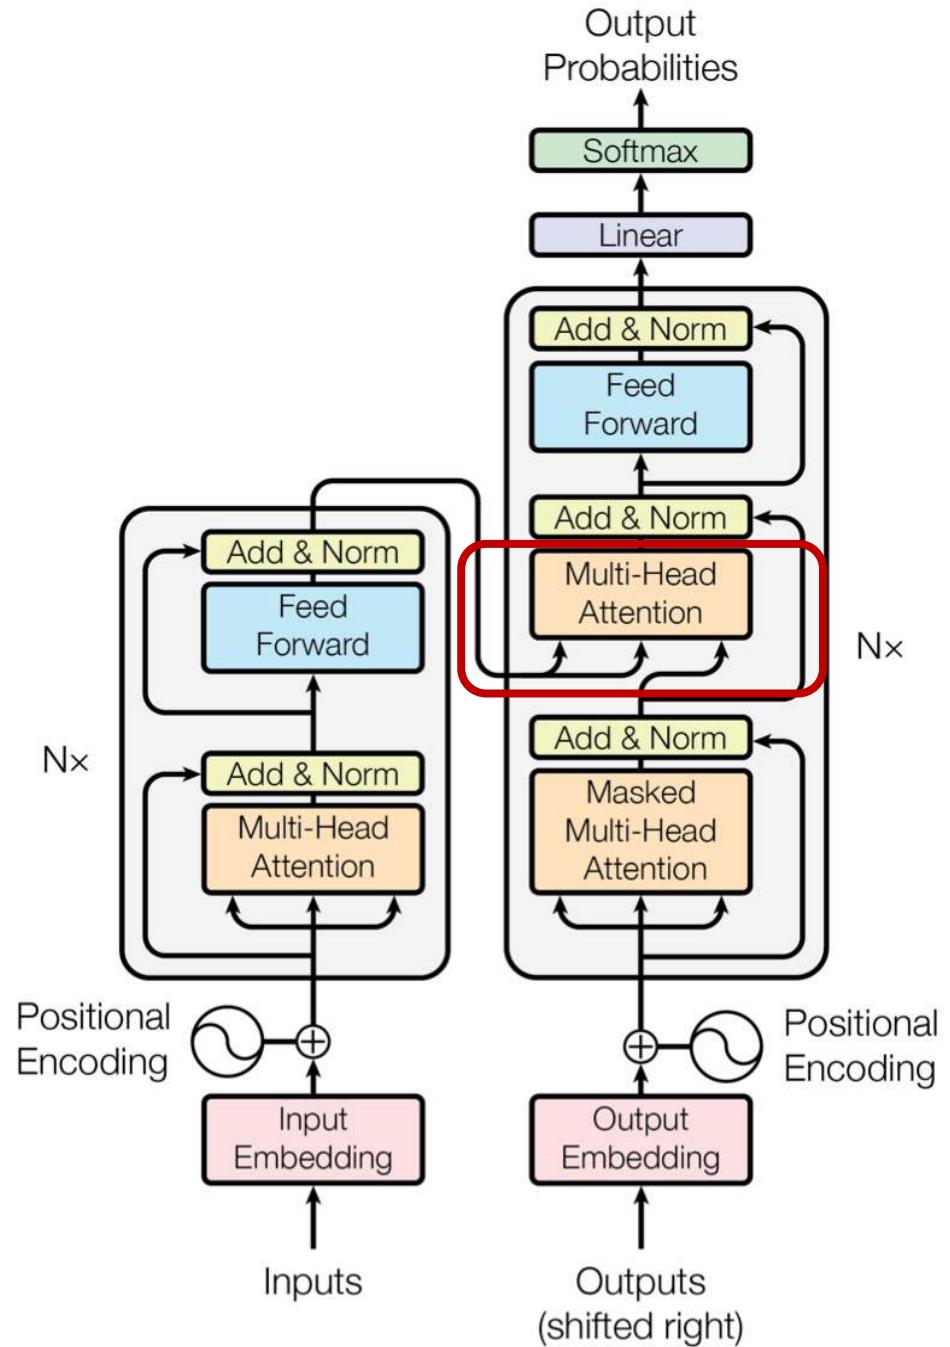
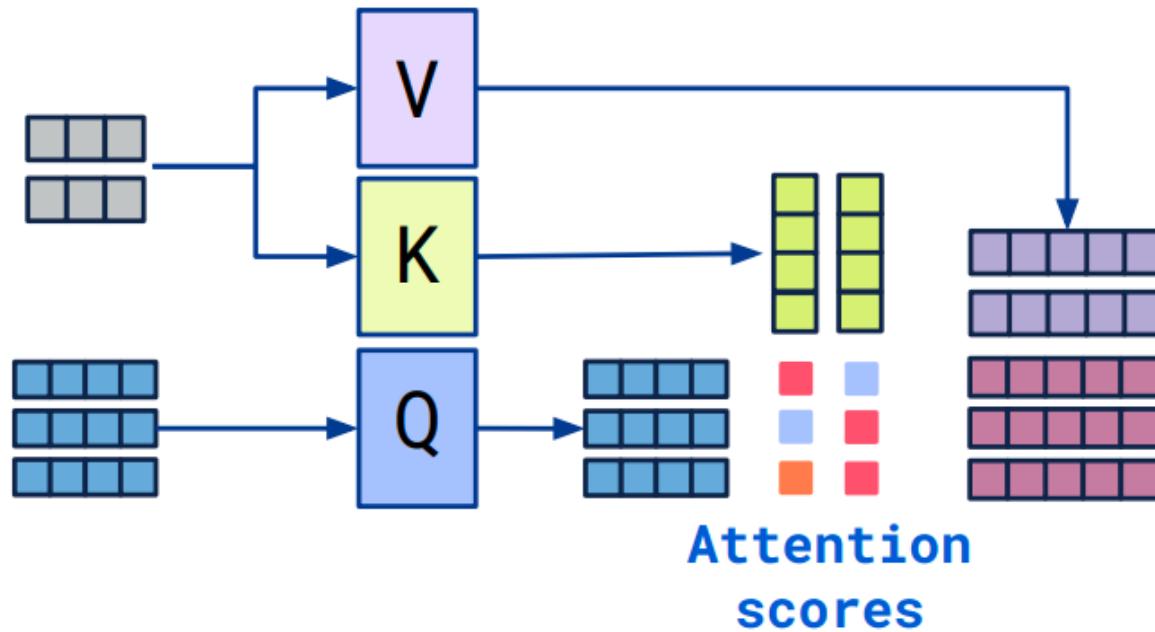
BERT: Encoder-only

- Predict a random 15% of (sub)word tokens.
 - Replace input word with [MASK] 80% of the time [Predict these!]
 - Replace input word with a random token 10% of the time
 - Leave input word unchanged 10% of the time (but still predict it!)
- The performance of BERT is better than GPT on lots of tasks. But it is not suitable for generation!



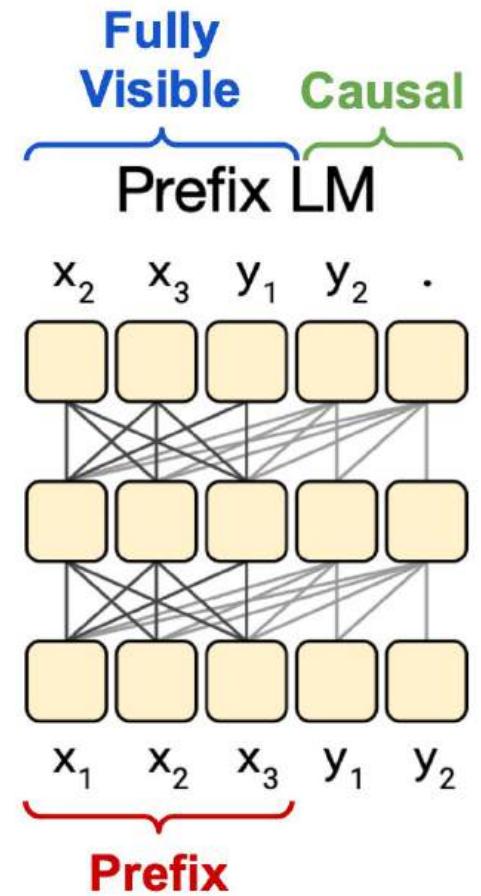
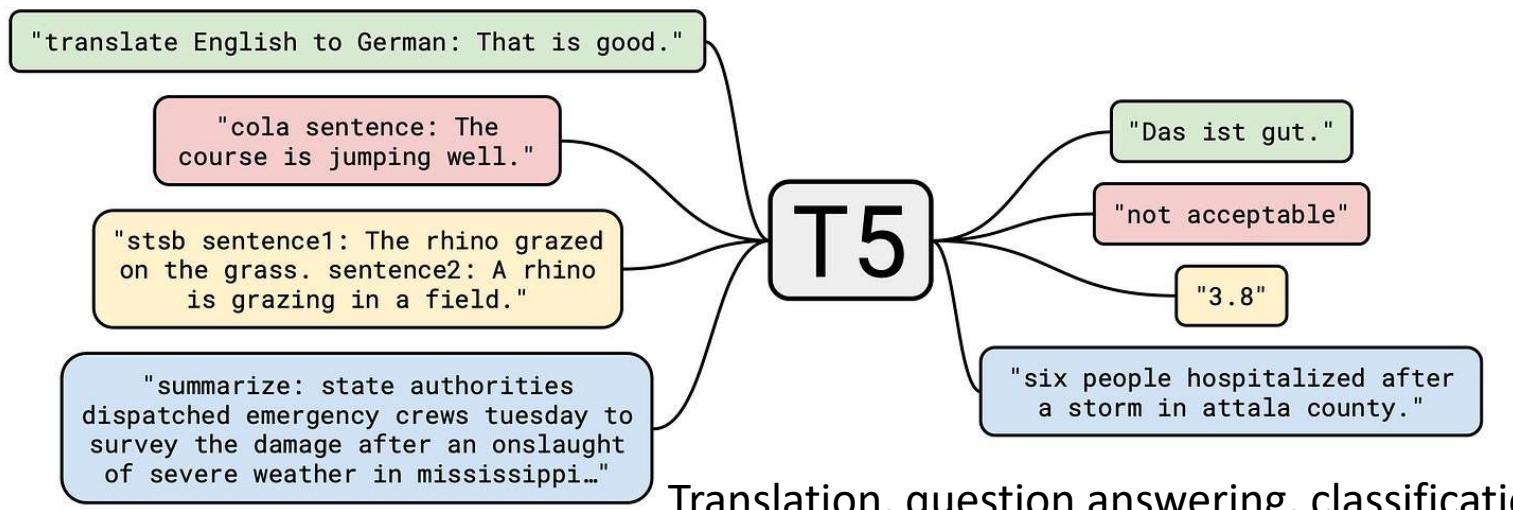
Encoder-decoder

- Suitable for seq-to-seq tasks
- **Cross attention:** embed encoder outputs into decoder



Text-to-Text Framework

- **T5 (The Unified Text-to-Text Transformer)**: model and solve a variety of different tasks with a shared approach: *the same model, objective, training procedure and decoding process*
- Encoder–decoder models are excellent for conditional generation ,but less efficient for free-form generation.



Post-training

1. Supervised Fine-tuning (GPT1)
2. Emergent Zero-shot Learning via Prompting (GPT2)
3. In-context (Few-shot) Learning (GPT3)
4. Chain-of-thought Prompting
5. Reinforcement Learning from Human Feedback ([covered in later class](#))

Supervised Fine Tuning (SFT)

- **Supervised fine-tuning:** continuing the training of a pretrained language model using **labeled data** for a **specific task**.
- Input–output pairs are often formatted as **prompts** and responses.
- Even a few thousand examples can make a big difference.

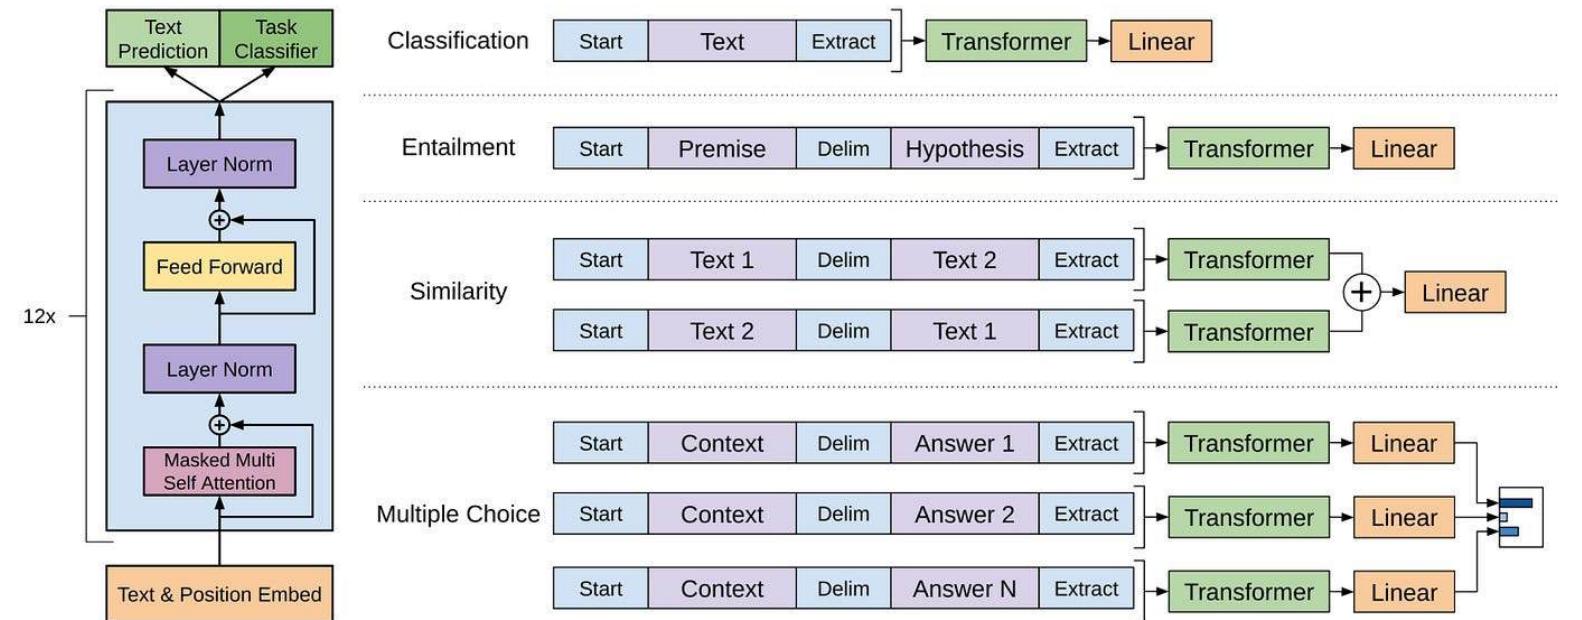


Figure 1: (left) Transformer architecture and training objectives used in this work. (right) Input transformations for fine-tuning on different tasks. We convert all structured inputs into token sequences to be processed by our pre-trained model, followed by a linear+softmax layer.

Zero-shot Learning

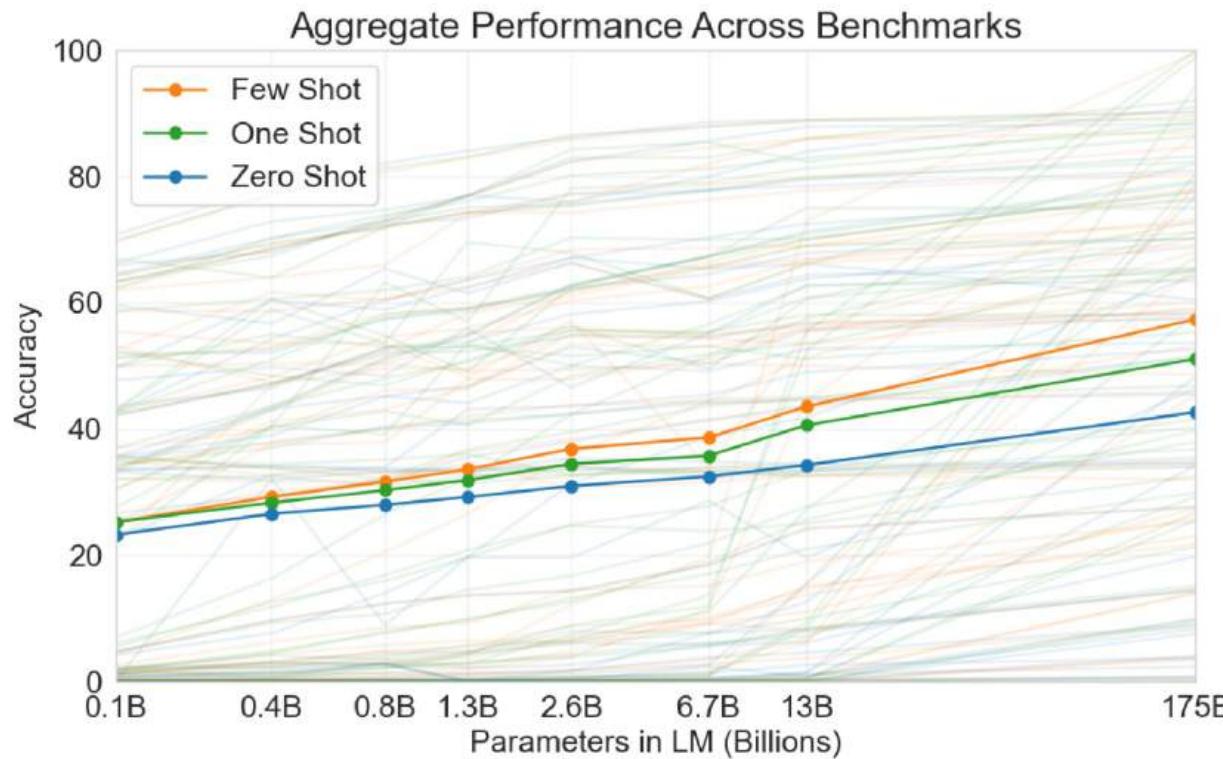
- GPT2: Language Models are Unsupervised Multitask Learners, 2019
 - Same Transformer decoder architecture as GPT-1, just larger 117 M → 1.5 B parameters
 - Trained on much more diverse web data: 4 GB → 40 GB of internet text (WebText)
- GPT-2 could **perform unseen tasks** directly from a prompt

Translate English to French: cat → ?

- GPT2 beats SOTA on some tasks without fine-tuning.

Few-shot Learning

- GPT3: Language Models are Few-Shot Learners



Translate English to French:
cat → ?

Translate English to French:
cat → chat
house → ?

Translate English to French:
cat → chat
dog → chien
house → ?

In-Context Learning

- In-Context Learning (ICL): learning from a few examples in the context

Input: 2014-06-01

Output: !06!01!2014!

Input: 2007-12-13

Output: !12!13!2007!

Input: 2010-09-23

Output: !09!23!2010!

Input: 2005-07-23

Output: !07!23!2005!

*in-context
examples*

test example

It allows models to **learn a task on the fly**, simply by being shown examples **within the prompt**, without any weight updates.



— — —

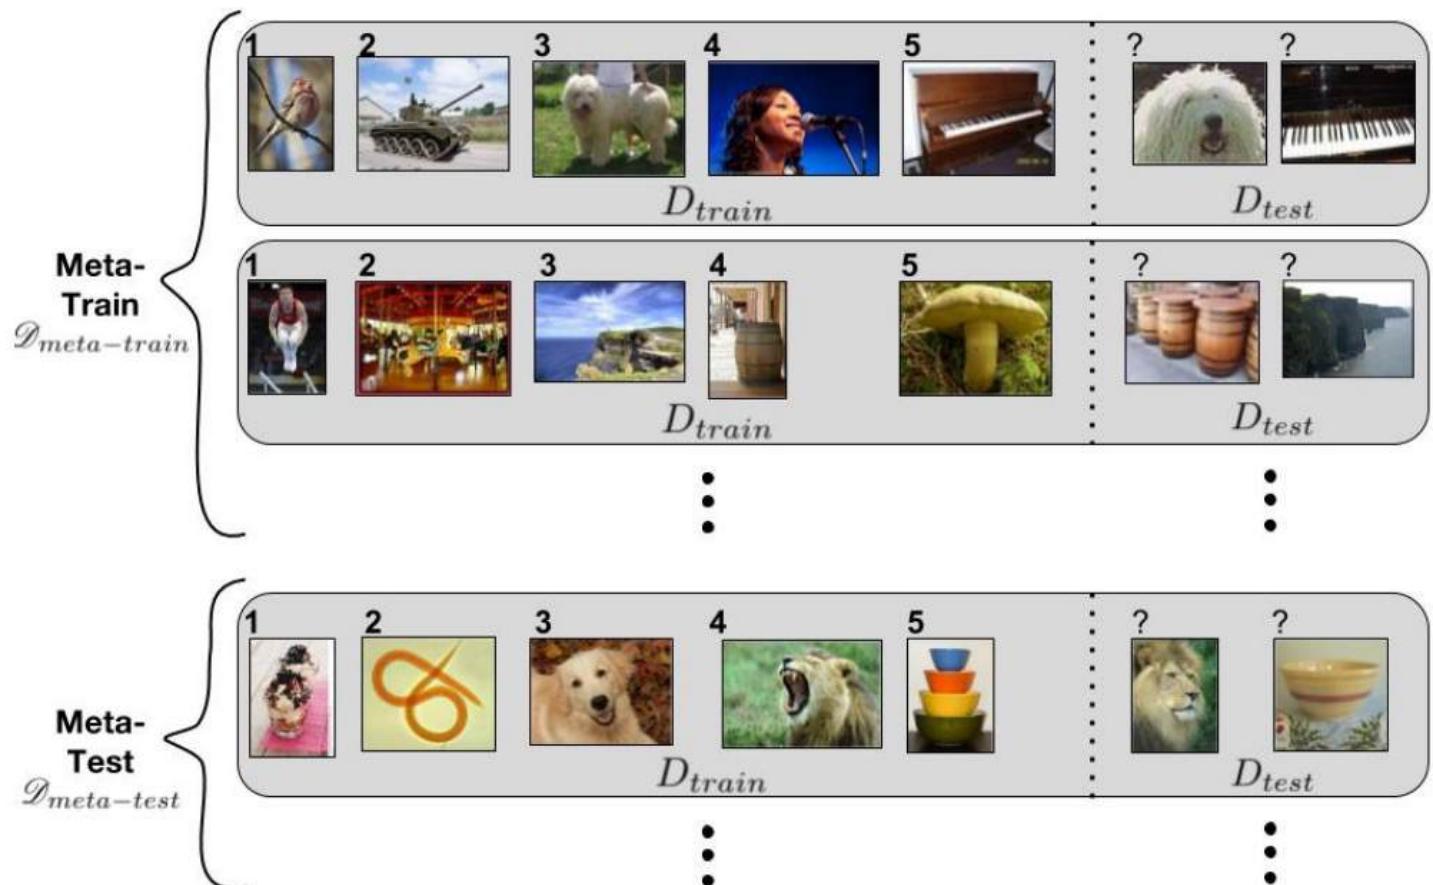
model completion

In-Context Learning

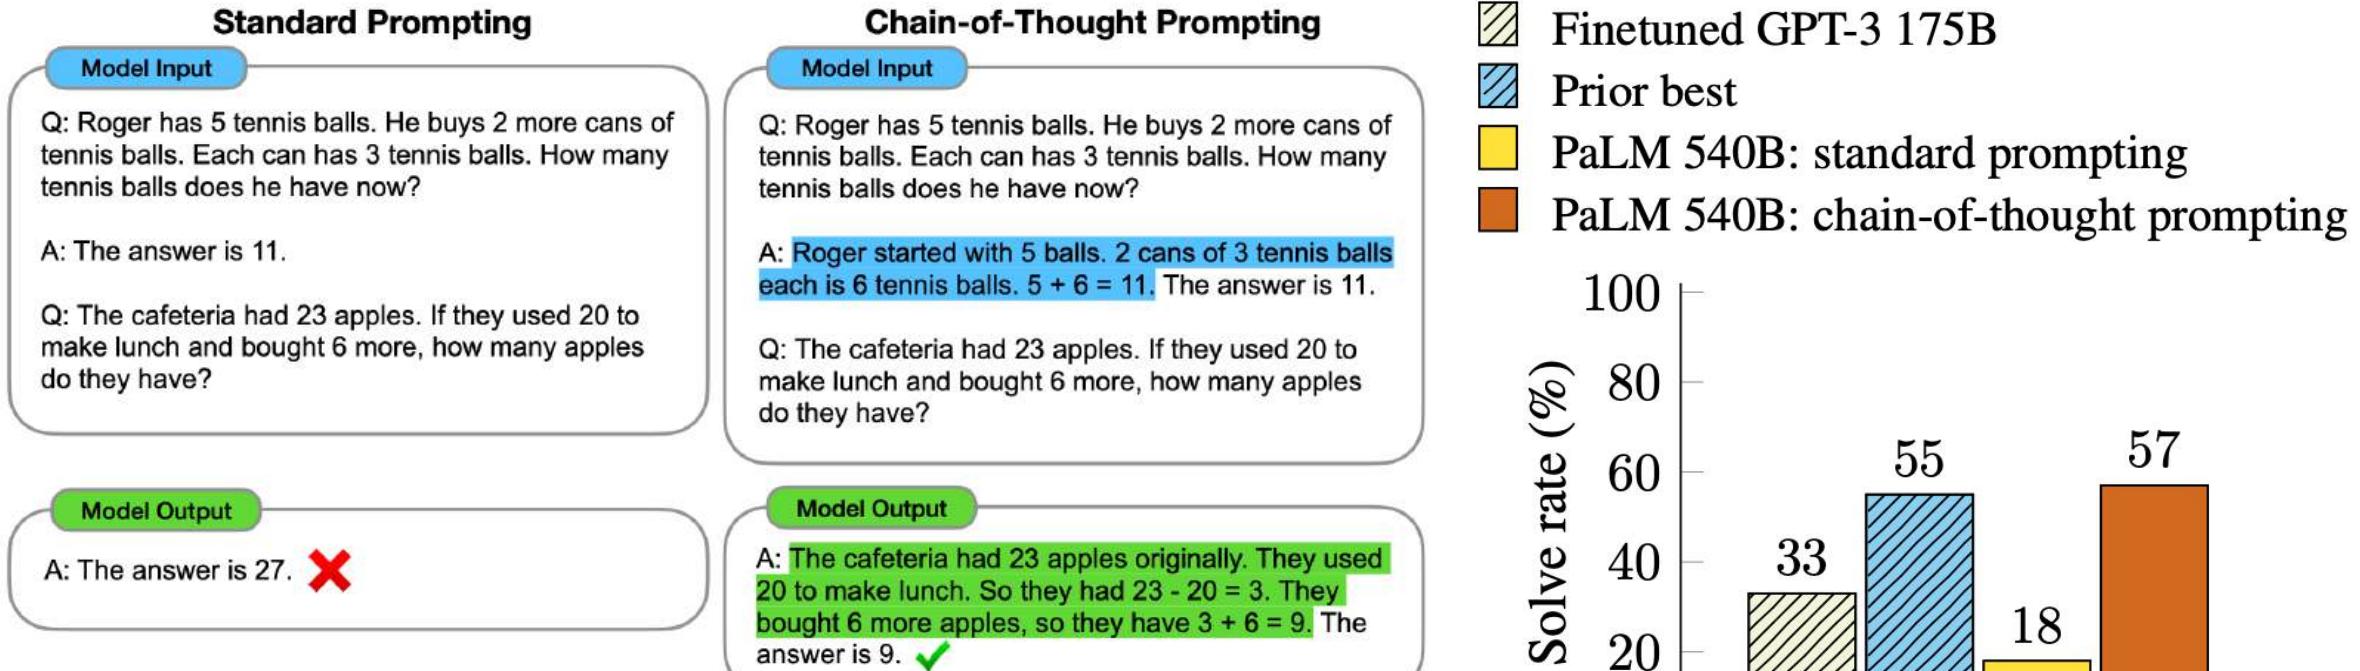
- Why it works?
 - It emerges in large-scale pretraining (**Emergent Abilities**)
 - Exposure to many formats and patterns in data (**SFT**)
 - The model's implicit **meta-learning** behavior: it acts like it's learning a task from examples, even though it's just predicting text
- But no long-term memory — model forgets after each prompt

Meta Learning

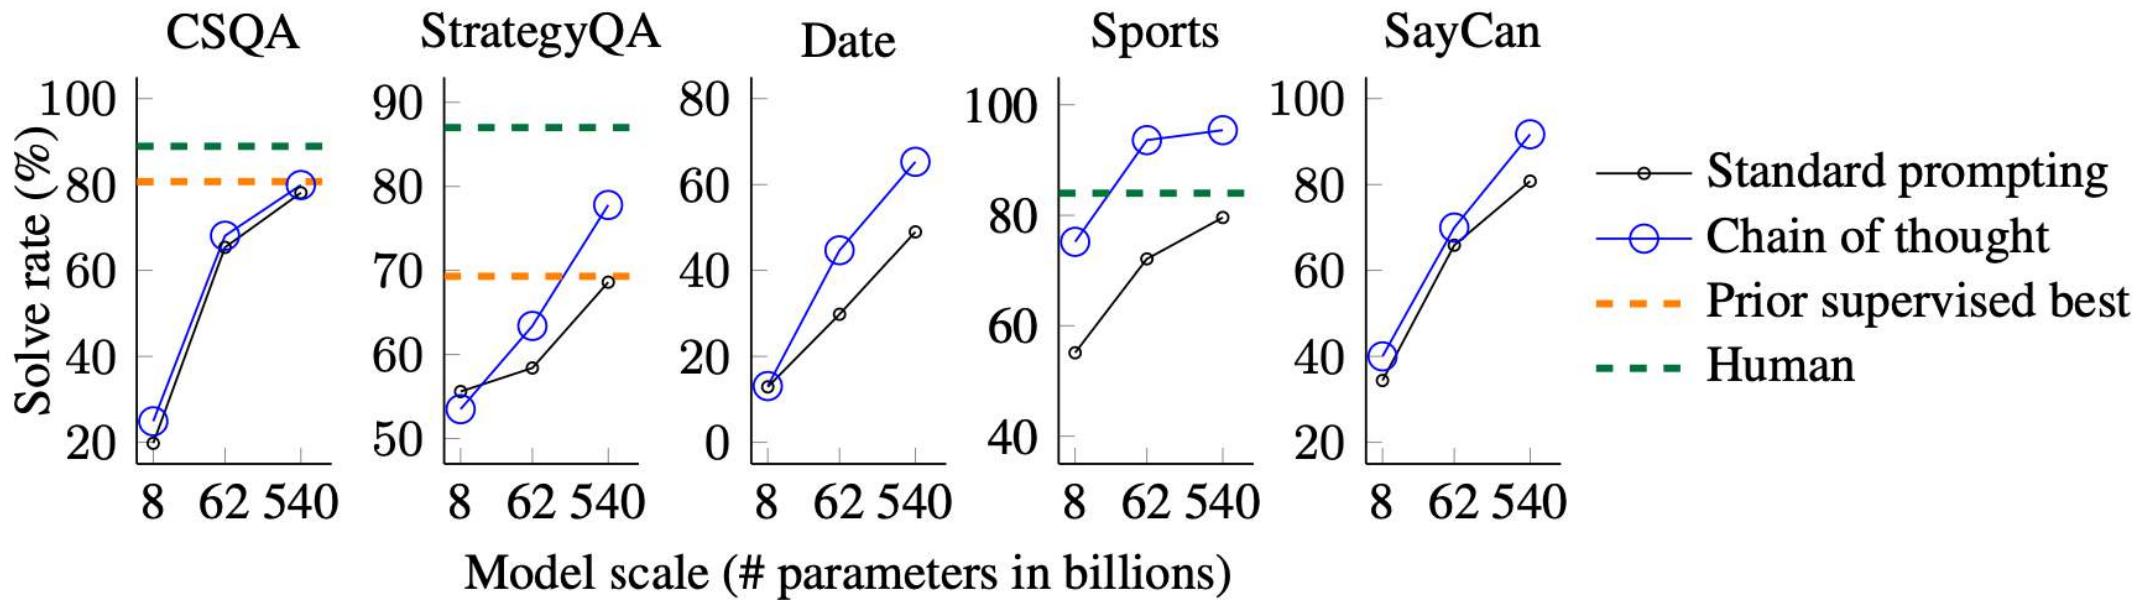
- Training Data
 - Inputs: $\{D_{tr}, x_{test}\}$
 - Targets: y_{test}
- Testing Data
 - Input: $\{D_{tr}^0, x_{test}^0\}$
 - Expected Output: y_{test}^0
- Meta Goal: learn how to learn



Chain-of-thought Promoting



CoT: Emergency



Common sense reasoning tasks

Sampling Strategies

How can we sample from $p_\theta(x_i|x_1, \dots, x_{i-1})$ at each step?

Option 1: Greedy Decoding

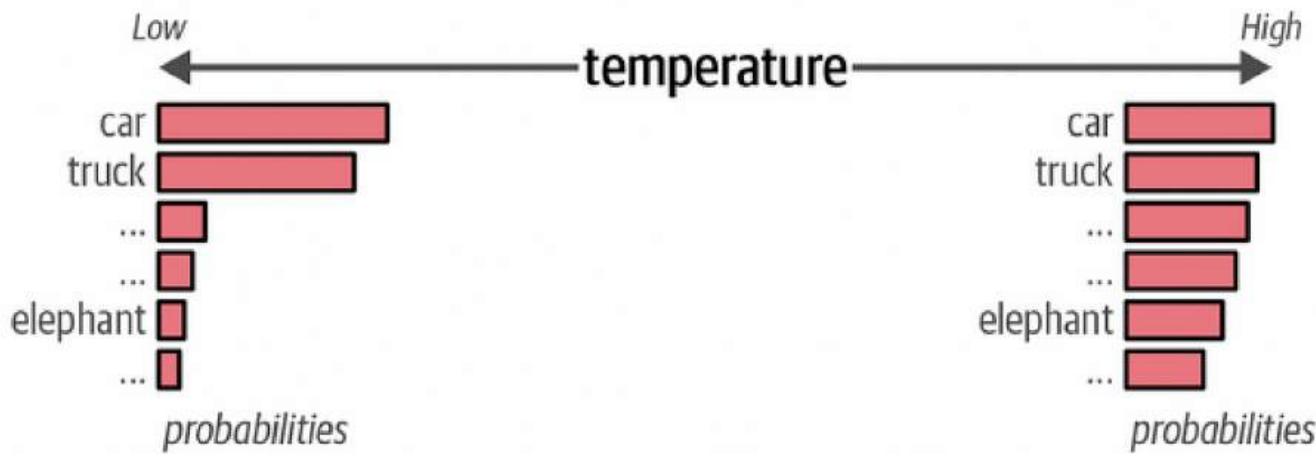
Pick the most likely token: $\hat{x}_i = \arg \max p_\theta(x_i|x_1, \dots, x_{i-1})$

- Pros: Deterministic, fast
- Cons: Low diversity, prone to repetition, gets stuck in loops

Sampling Strategies

Option 2: Temperature Sampling: Adjusting the Creativity

- High Temperature (≥ 1): Creates more varied and creative outputs. Ideal for brainstorming or fiction writing.
- Low Temperature (~ 0): Produces more focused, predictable responses, looks like argmax, ideal for tasks like text summarization or factual Q&A.



v_k : the k-th vocabulary token,
 l_k : the token's logit

$$p(v_k) = \frac{e^{\frac{l_k}{T}}}{\sum_k e^{\frac{l_k}{T}}}$$

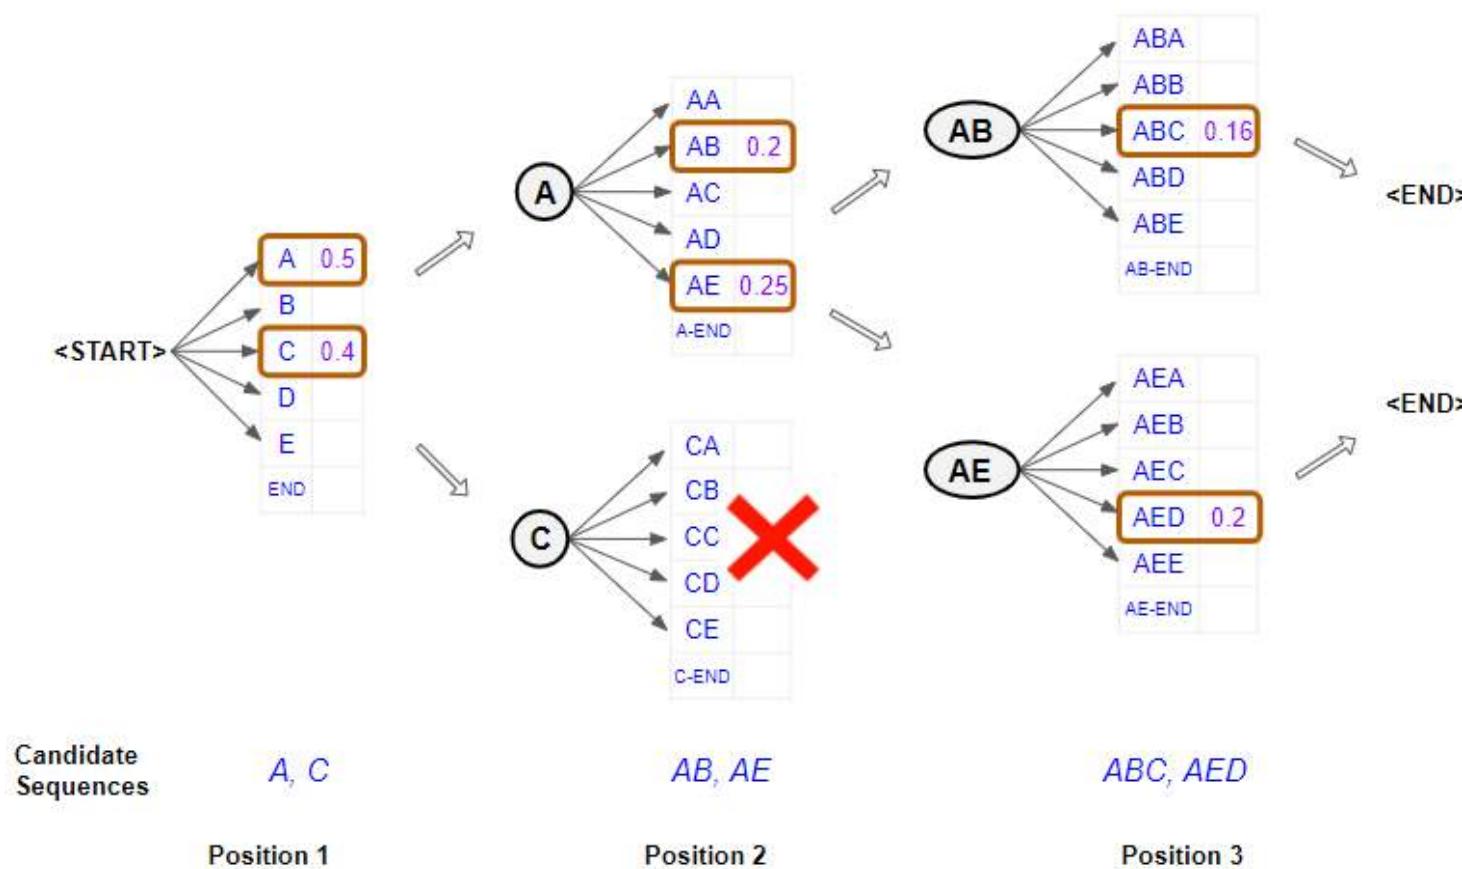
Sampling Strategies

Option 3: Top-K sampling limits the model's options to the top K most likely words

Option 4: Nucleus Sampling introduces sparsity, choose smallest set of tokens whose cumulative probability $\geq p$

Option 5: Beam search maintains the top B most likely partial sequences (called beams) at each decoding step.

Beam Search



- Path probability

$$p_{\theta}(x_1, \dots, x_l) = \prod_{i=1}^l p_{\theta}(x_i | x_1, \dots, x_{i-1})$$

- Longer sequences naturally have smaller probabilities due to multiplying more terms.
- Fix: use length normalization:

$$\text{score} = \frac{1}{(5 + |\text{seq}|)^a / (5 + 1)^a} \cdot \log P(\text{seq})$$

PixelRNN

- Decompose an image to pixels

$$p_{model}(x) = \prod_{i=1}^N p_{model}(x_i|x_1, \dots, x_{i-1})$$

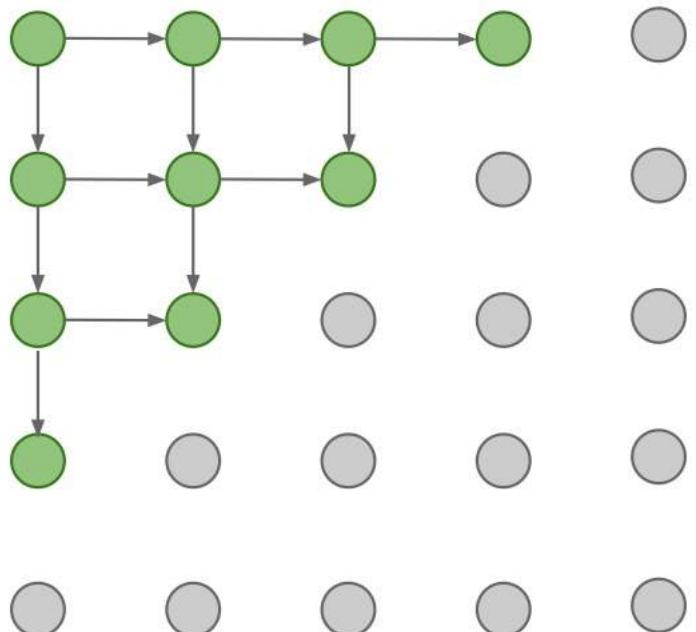
- Loss:

$$\sum_{x \in D} \log p_{model}(x|\theta) = \sum_{x \in D} \sum_i^N \log p_{model}(x_i|x_1, \dots, x_{i-1})$$

- $p_{model}(x_i|x_1, \dots, x_{i-1})$: discrete probability on [0,1,...,255]

PixelRNN

- Generate image pixels starting from corner



Generation: compute $p_{model}(x_i|x_1, \dots, x_{i-1}) \in \mathbb{R}^{256}$,
then take x_i as the value with the highest probability.

The sample generation process is
extremely slow for high-resolution
images!

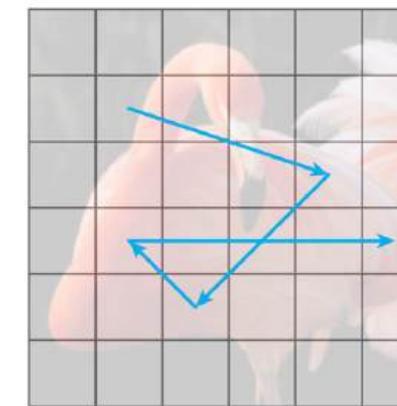
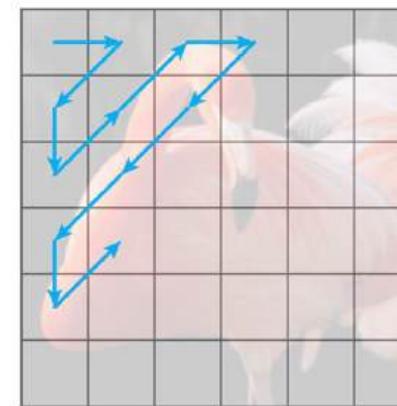
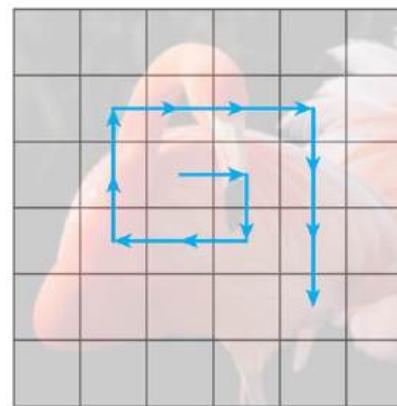
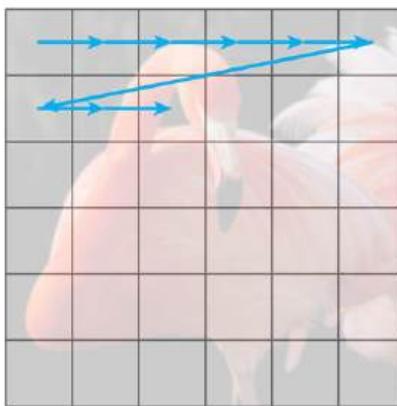
Recaps: Joint Distribution Decomposition

- A joint distribution can be written as a product of conditionals **in any order**

$$\begin{aligned} p(A, B, C) &= p(A)p(B \mid A)p(C \mid A, B) \\ &= p(A)p(C \mid A)p(B \mid A, C) \\ &= p(B)p(A \mid B)p(C \mid A, B) \\ &= p(B)p(C \mid B)p(A \mid B, C) \\ &= p(C)p(A \mid C)p(B \mid A, C) \\ &= p(C)p(B \mid C)p(A \mid B, C) \end{aligned}$$

Dependency Order

- The ordering may effect the training complexity



Joint Distribution Decomposition

- A joint distribution can be written as a product of conditionals **in any partition**

$$\begin{aligned} p(A, B, C, D) &= p(A, B)p(C, D \mid A, B) \\ &= p(C, D)p(A, B \mid C, D) \\ &= p(A, B, C)p(D \mid A, B, C) \\ &= \dots \end{aligned}$$

Joint Distribution Decomposition

- Partitioning in the input space

$$p(x_1, x_2, \dots, x_n) = p(x_1)p(x_2 \mid x_1)\dots p(x_n \mid x_1, x_2, \dots, x_{n-1})$$

- Partition in the latent space (e.g. VAE to come next)

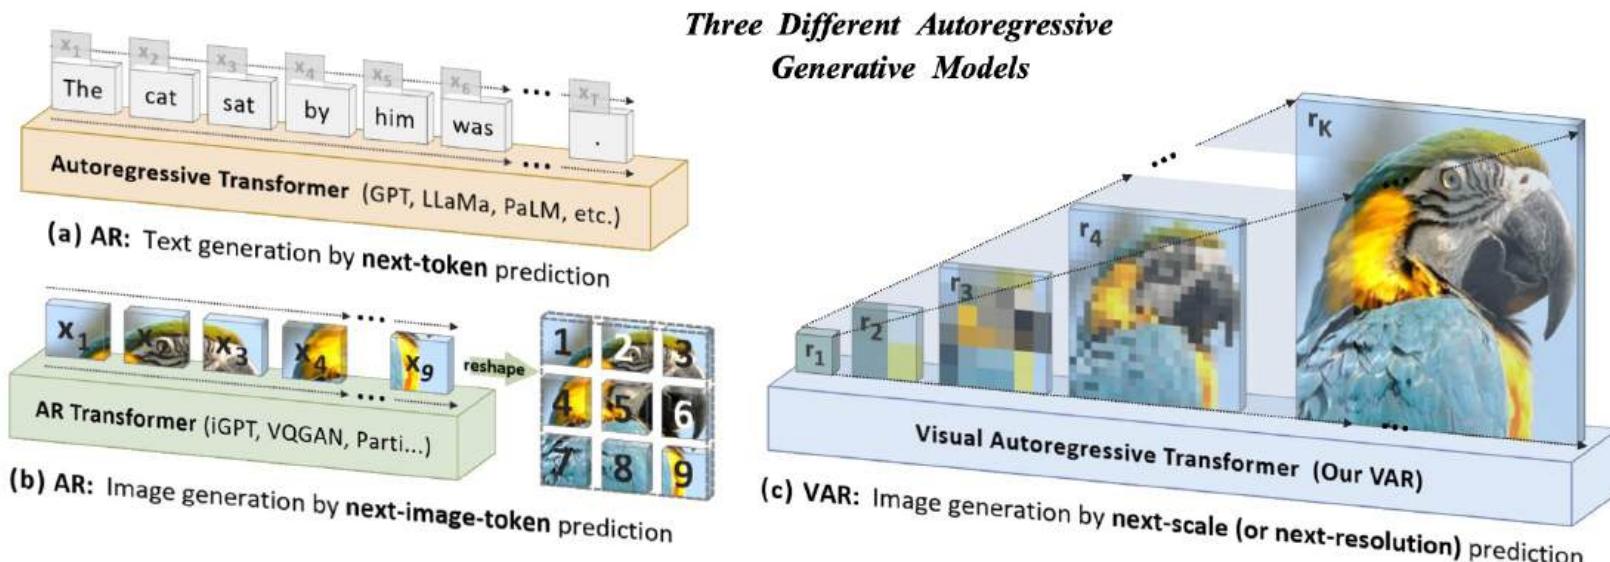
$$p(\mathbf{x}, \mathbf{z}) = p(\mathbf{z})p(\mathbf{x} \mid \mathbf{z})$$

with $p(\mathbf{z}) = p(z_1)p(z_2 \mid z_1)\dots p(z_n \mid z_1, z_2, \dots, z_{n-1})$

Visual Autoregressive Modeling

- Next-scale Prediction: (r_1, r_2, \dots, r_k) multi scale representation

$$p(r_1, r_2, \dots, r_K) = \prod_{k=1}^K p(r_k | r_1, r_2, \dots, r_{k-1}),$$



Next-token Prediction

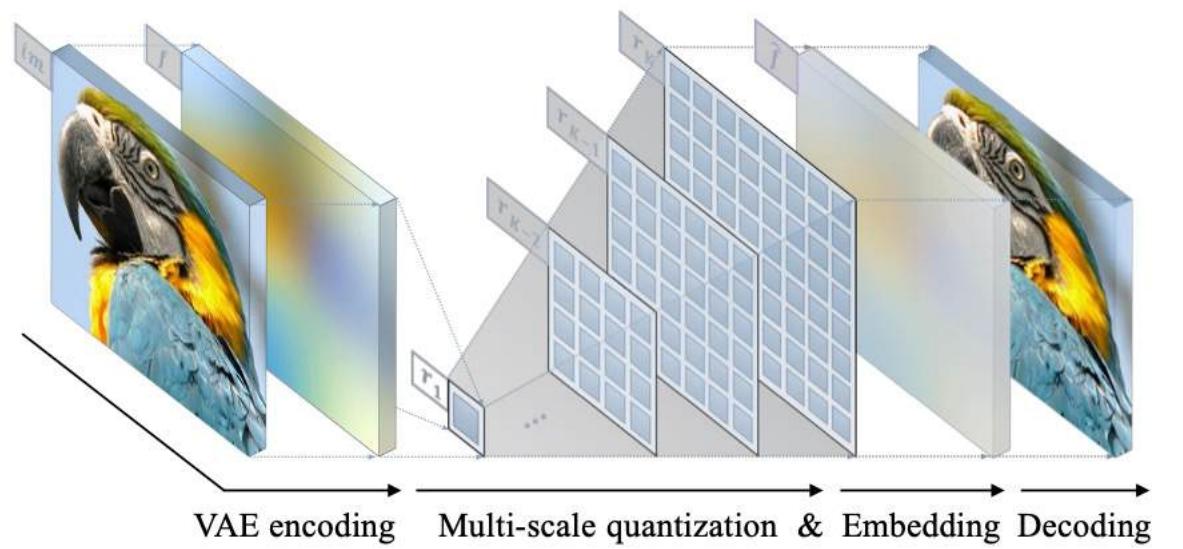
- A quantized autoencoder convert the image feature map to discrete tokens
- Encoder

$$f = \mathcal{E}(im), \quad q = \mathcal{Q}(f), \\ q^{(i,j)} = \left(\arg \min_{v \in [V]} \| \text{lookup}(Z, v) - f^{(i,j)} \|_2 \right) \in [V],$$

- Decoder

$$\hat{f} = \text{lookup}(Z, q), \quad \hat{im} = \mathcal{D}(\hat{f}),$$

Stage 1: Training multi-scale VQVAE on images
(to provide the ground truth for training Stage 2)



Auto regressive learning in latent code space

Next-token Prediction

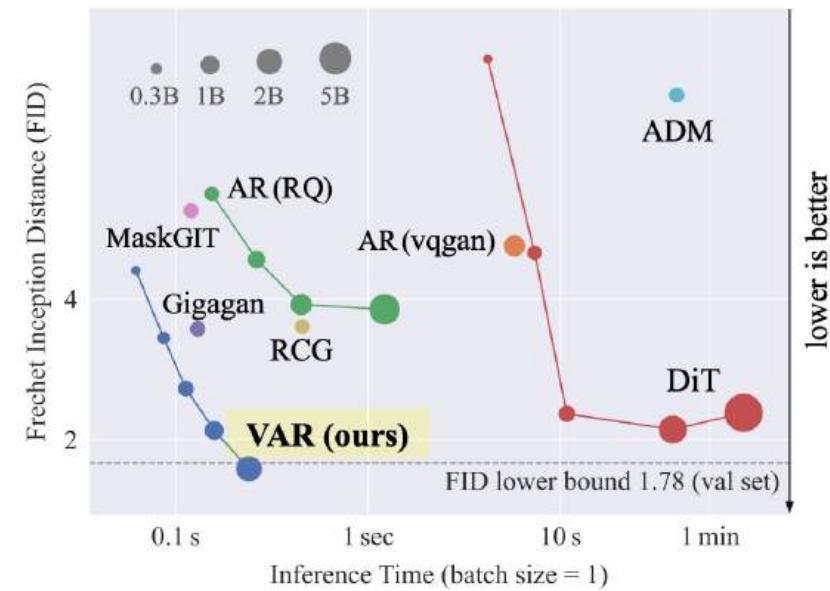
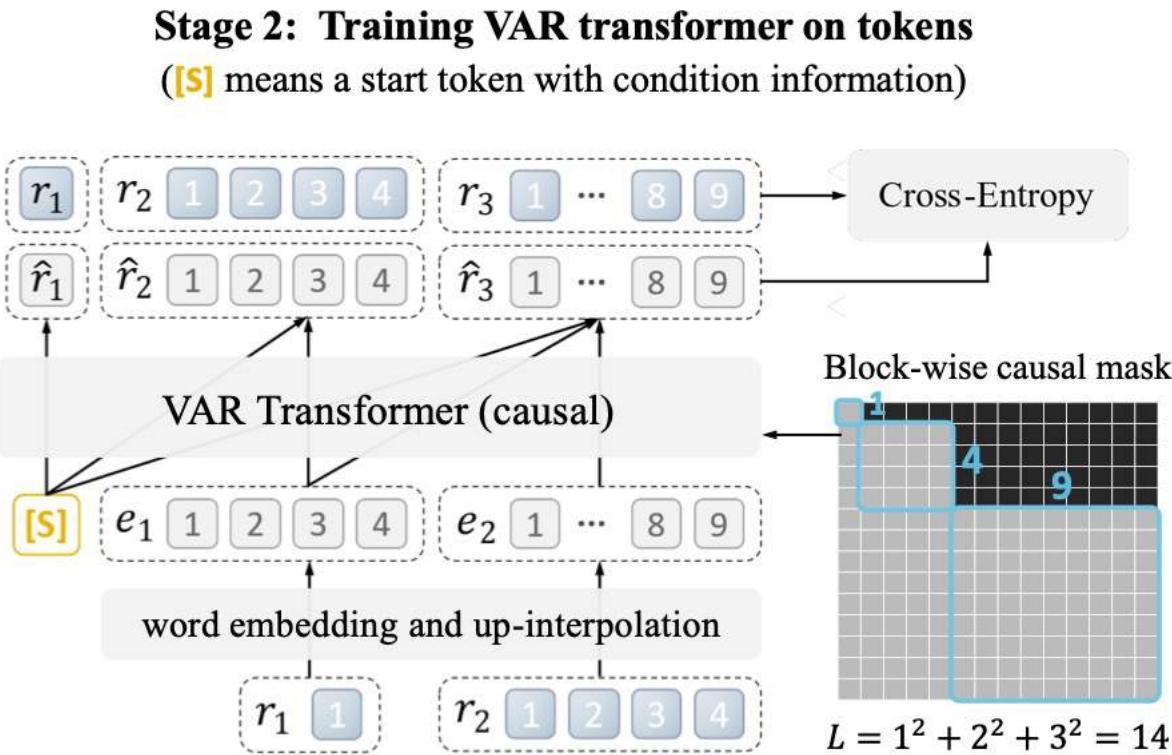


Figure 3: Scaling behavior of different model families on ImageNet 256×256 generation benchmark. The FID of the validation set serves as a reference lower bound (1.78). VAR with 2B parameters reaches an FID of 1.73, surpassing L-DiT with 3B or 7B parameters.

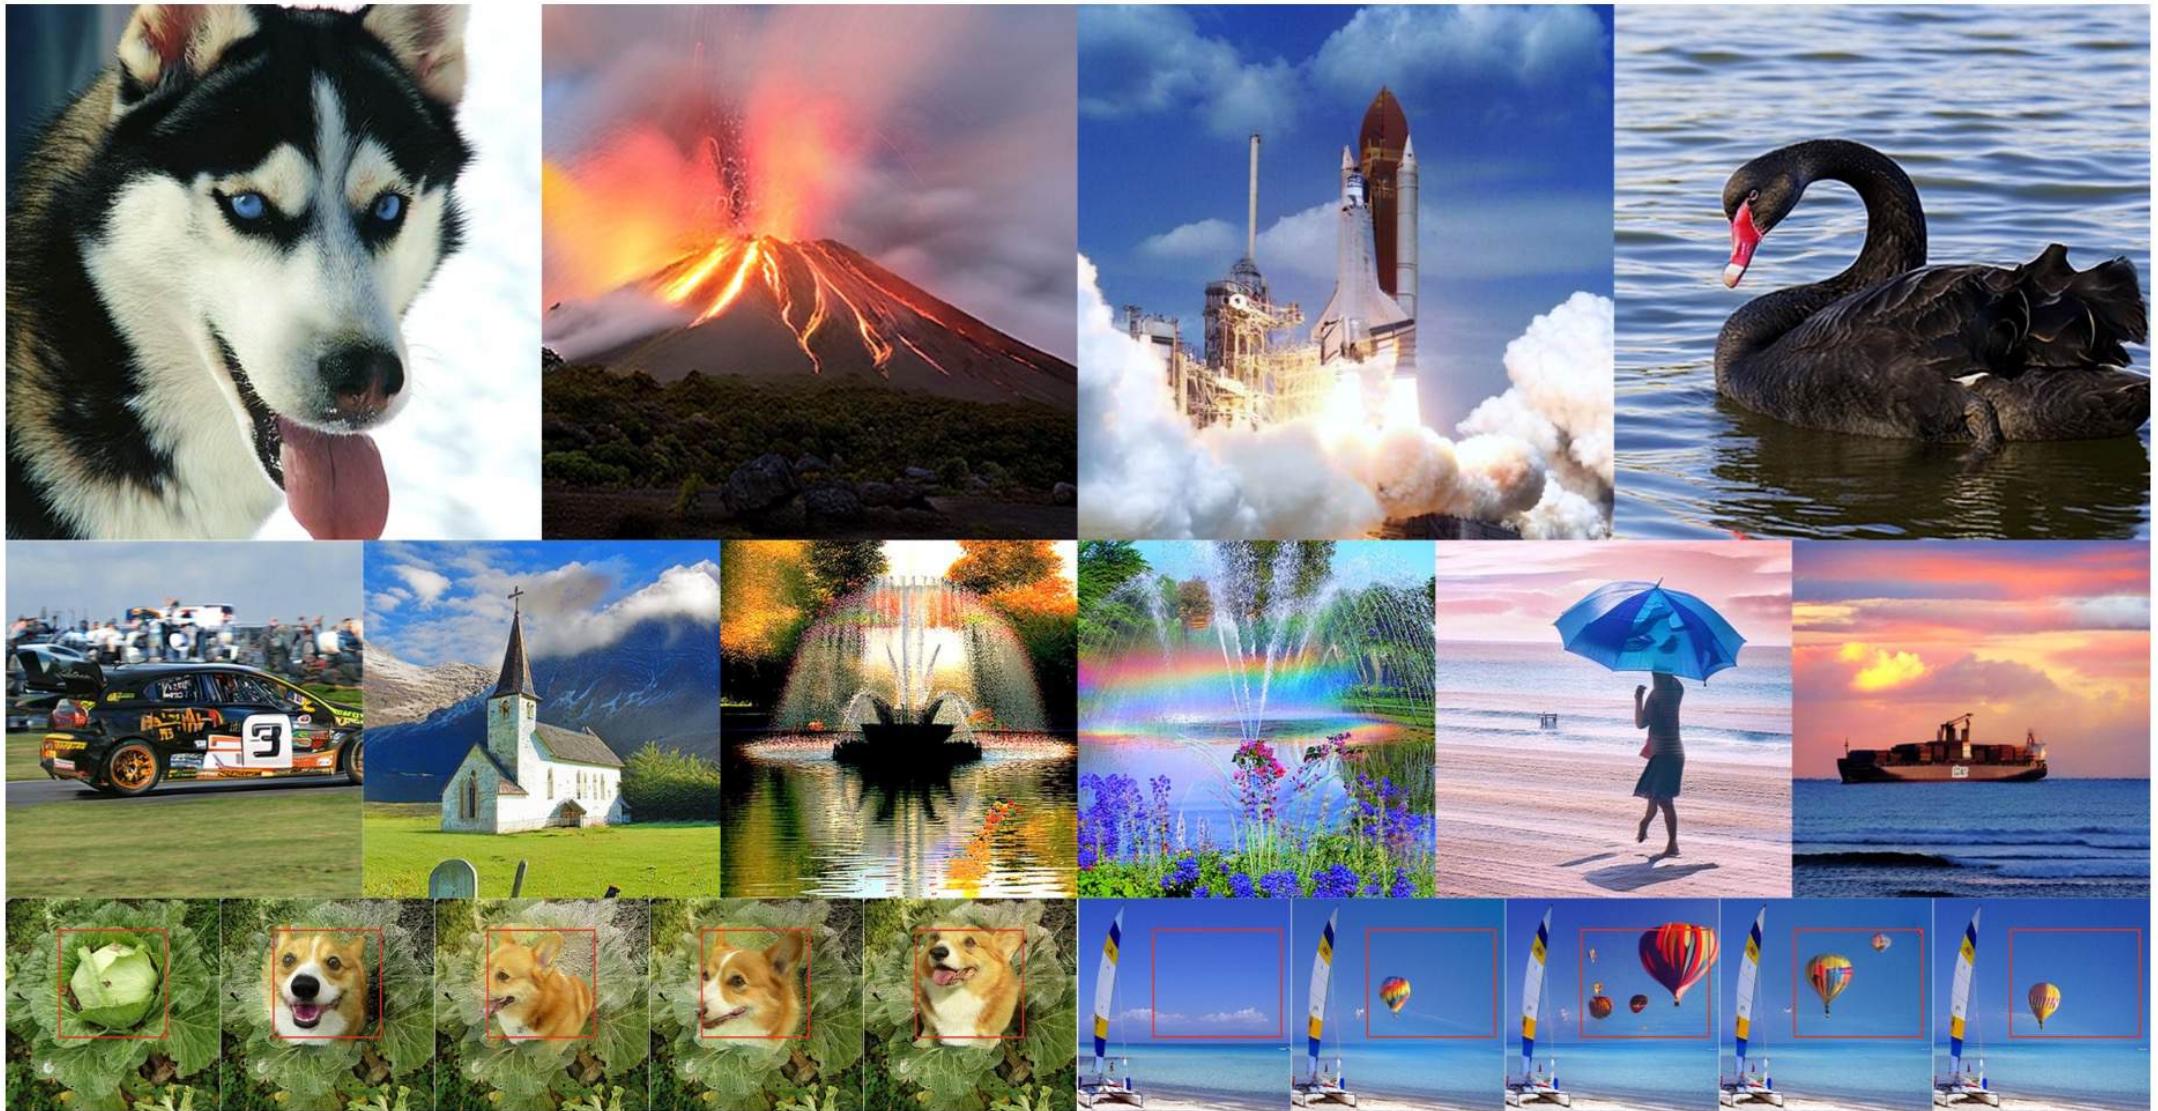
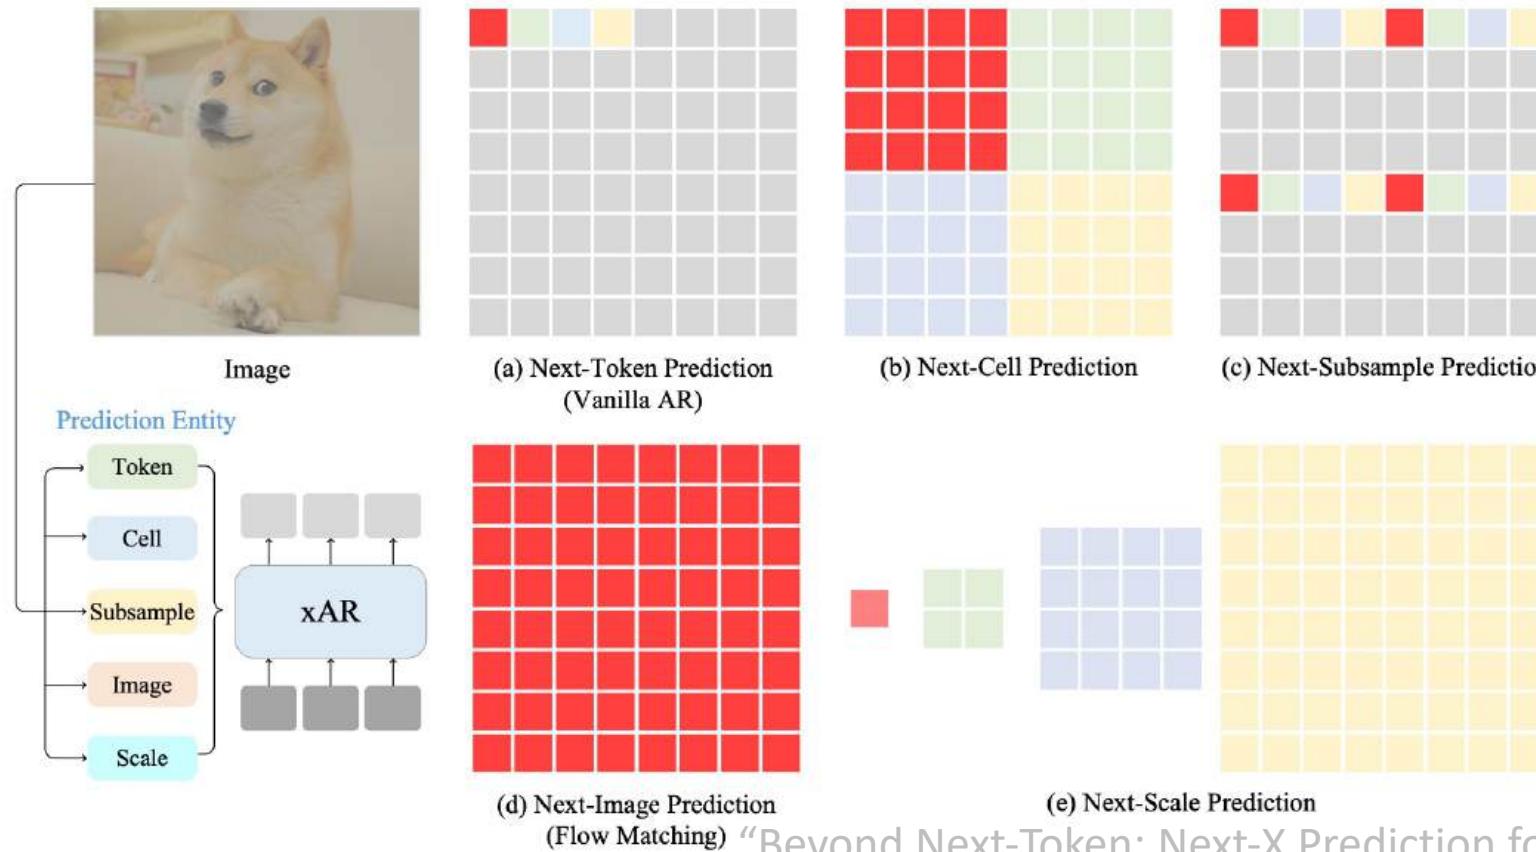


Figure 1: **Generated samples from Visual AutoRegressive (VAR) transformers trained on ImageNet.** We show 512×512 samples (top), 256×256 samples (middle), and zero-shot image editing results (bottom).

Visual Autoregressive Modeling

- Next-X Prediction: X can be any representation



"Beyond Next-Token: Next-X Prediction for Autoregressive Visual Generation"